

Course: Introduction to AI

Prof. Ahmed Guessoum

The National Higher School of AI

Chapter 6

Constraint Satisfaction Problems

Outline

- Defining Constraint Satisfaction Problems
 - ♦ Example problem: Map colouring
 - ♦ Example problem: Job-shop scheduling
 - ♦ Variations on the CSP formalism
- Constraint Propagation: Inference in CSPs
 - ♦ Node consistency
 - ♦ Arc consistency
 - ♦ Path consistency
 - ♦ K -consistency
 - ♦ Global constraints
 - ♦ Sudoku example

Outline

- Backtracking Search for CSPs
 - ◆ Variable and value ordering
 - ◆ Interleaving search and inference
 - ◆ Intelligent backtracking: Looking backward
- Local Search for CSPs
- The Structure of Problems

Constraint Satisfaction Problems

- Chapters 3 and 4: problems can be solved by searching in a space of **states**.
- These states can be evaluated by domain-specific heuristics and tested to see whether they are goal states.
- The search algorithm “sees” each state as atomic, or indivisible—a black box with no internal structure.
- In this chapter, each state is seen as a set of variables, each of which has a value.
- A problem is solved when each variable has a value that satisfies all the constraints on the variable.
- A problem described this way is called a **Constraint Satisfaction Problem**, or CSP.
- CSP search algorithms use general-purpose rather than problem-specific heuristics to solve complex problems.

Defining CSPs

- A constraint satisfaction problem consists of three components (sets): X , D , and C :
 - ♦ X is a set of variables, $\{X_1, \dots, X_n\}$.
 - ♦ D is a set of domains, $\{D_1, \dots, D_n\}$, one for each variable.
 - Each domain D_i consists of a set of allowable values, $\{v_1, \dots, v_k\}$ for variable X_i .
 - ♦ C is a set of constraints that specify allowable combinations of values.
 - Each constraint C_i consists of a pair $\langle \textit{scope}, \textit{rel} \rangle$, where *scope* is a tuple of variables that participate in the constraint and *rel* is a relation that defines the values that those variables can take.

Defining CSPs

- A relation can be represented as:
 - ♦ an explicit list of all tuples of values that satisfy the constraint, **or**
 - ♦ an abstract relation that supports two operations:
 - testing if a tuple is a member of the relation and
 - enumerating the members of the relation.
- E.g., if X_1 and X_2 both have the domain $\{A,B\}$, then the constraint that they must have different values can be written as $\langle (X_1, X_2), [(A,B), (B,A)] \rangle$ or as $\langle (X_1, X_2), X_1 \neq X_2 \rangle$.

Defining CSPs

- To solve a CSP, we need to define a state space and the notion of a solution.
- A CSP state is defined by an **assignment** of values to **some or all** of the variables, $\{X_i = v_i, X_j = v_j, \dots\}$.
- An assignment that does not violate any constraints is called a **consistent** or legal assignment.
- A **complete assignment** is one in which every variable is assigned.
- A **solution** to a CSP is *a consistent, complete assignment*.
- A **partial assignment** is one that assigns values to only some of the variables.

Example problem: Map colouring




Figure 6.1 (a) The principal states and territories of Australia. Coloring this map can be viewed as a constraint satisfaction problem (CSP). The goal is to assign colors to each region so that no neighboring regions have the same color. (b) The map-coloring problem represented as a constraint graph.

Map colouring problem

- **Task:** colour each region of the map either *Red*, *Green*, or *Blue* so that no neighbouring regions have the same colour.
- CSP formulation:
 - ♦ The variables are the regions
 $X = \{WA, NT, Q, NSW, V, SA, T\}$.
 - ♦ Domain of each variable: the set
 $D_i = \{\text{red}, \text{green}, \text{blue}\}$.
 - ♦ There are nine constraints:
 $C = \{SA \neq WA, SA \neq NT, SA \neq Q, SA \neq NSW, SA \neq V, WA \neq NT, NT \neq Q, Q \neq NSW, NSW \neq V\}$
- N.B.: $SA \neq WA$ can be fully enumerated in turn as
 $\{(\text{red}, \text{green}), (\text{red}, \text{blue}), (\text{green}, \text{red}), (\text{green}, \text{blue}), (\text{blue}, \text{red}), (\text{blue}, \text{green})\}$

Map colouring problem

- A CSP can be visualized as a **constraint graph** (see previous figure).
 - ♦ The graph nodes correspond to problem variables.
 - ♦ A link connects any two variables that participate in a constraint.
- Australia map colouring problem: many possible solutions to this. One possible solution is:
 $\{WA=\text{red}, NT=\text{green}, Q=\text{red}, NSW=\text{green}, V=\text{red}, SA=\text{blue}, T=\text{red}\}.$

Example problem: Map colouring




Figure 6.1 (a) The principal states and territories of Australia. Coloring this map can be viewed as a constraint satisfaction problem (CSP). The goal is to assign colors to each region so that no neighboring regions have the same color. (b) The map-coloring problem represented as a constraint graph.

Problem formulation as a CSP

- Why formulate a problem as a CSP?
 - ♦ CSPs have a natural representation for a wide variety of problems.
 - ♦ If a CSP-solving system is available, it is often easier to solve a problem using it than design a custom solution using another search technique.
 - ♦ CSP solvers can be faster than state-space searchers because they can quickly eliminate large parts of the search space. E.g. if $\{SA=blue\}$
 - $\rightarrow$ none of the five neighbouring variables can take value blue. $\rightarrow$ for constraint propagation *blue* is never considered.
 - Other search techniques: consider $3^5 = 243$ assignments for the five neighbouring vars; $\rightarrow$
 - In CSP only $2^5 = 32$ assignments; a reduction of 87%.

Problem formulation as a CSP

- In regular state-space search: one can only ask if a specific state is a goal or not.
- With CSPs:
 - ♦ once a partial assignment is found not to be a solution, one can
 - immediately discard further refinements of the partial assignment.
 - see *why* the assignment is not a solution—i.e. which variables violate a constraint—so attention can be focused on the variables that matter.
- ➔ many problems that are intractable for regular state-space search can be solved quickly when formulated as a CSP. (See chapter on Planning)


Example problem: Job-shop scheduling

- Problem of scheduling the day's work (in terms of job/tasks) in a factory, subject to various constraints.
- Many of these problems are solved with CSP techniques.
- E.g. scheduling the assembly of a car.
 - ♦ The job is composed of tasks;
 - ♦ Each task can be modeled as a variable, whose value is the time the task starts, expressed as an integer number of minutes;
 - ♦ Constraints can state, for instance:
 - that one task must occur before another. E.g., a wheel must be installed before the hubcap is put on;
 - that only so many tasks can go on at once;
 - the amount of time it takes to complete.

Example: car assembly scheduling

- Suppose the problem consists of (only) 15 tasks: install axles (front and back), affix all four wheels (right and left, front and back), tighten nuts for each wheel, affix hubcaps, and inspect the final assembly.
- The tasks can be represented with 15 variables:

$X = \{\text{Axle}_F, \text{Axle}_B, \text{Wheel}_{RF}, \text{Wheel}_{LF}, \text{Wheel}_{RB}, \text{Wheel}_{LB}, \text{Nuts}_{RF}, \text{Nuts}_{LF}, \text{Nuts}_{RB}, \text{Nuts}_{LB}, \text{Cap}_{RF}, \text{Cap}_{LF}, \text{Cap}_{RB}, \text{Cap}_{LB}, \text{Inspect}\}$



Example: car assembly scheduling

- The value of each variable: the time the task starts.
- **Precedence constraints** between individual tasks:
 - ♦ Whenever a task T1 must occur before task T2, and task T1 takes duration d1 to complete, we add an arithmetic constraint: $T1 + d1 \leq T2$.
- In our example,
 - ♦ the axles must be placed before the wheels
 - ♦ it takes 10 minutes to install an axle. →
 $Axle_F + 10 \leq Wheel_{RF} ; \quad Axle_F + 10 \leq Wheel_{LF} ;$
 $Axle_B + 10 \leq Wheel_{RB} ; \quad Axle_B + 10 \leq Wheel_{LB} ;$

Example: car assembly scheduling

- ♦ For each wheel: affix it (1 minute), then tighten the nuts (2 minutes), and attach the hubcap (1 minute):

$$\text{Wheel}_{\text{RF}} + 1 \leq \text{Nuts}_{\text{RF}} ; \quad \text{Nuts}_{\text{RF}} + 2 \leq \text{Cap}_{\text{RF}} ;$$

$$\text{Wheel}_{\text{LF}} + 1 \leq \text{Nuts}_{\text{LF}} ; \quad \text{Nuts}_{\text{LF}} + 2 \leq \text{Cap}_{\text{LF}} ;$$

$$\text{Wheel}_{\text{RB}} + 1 \leq \text{Nuts}_{\text{RB}} ; \quad \text{Nuts}_{\text{RB}} + 2 \leq \text{Cap}_{\text{RB}} ;$$

$$\text{Wheel}_{\text{LB}} + 1 \leq \text{Nuts}_{\text{LB}} ; \quad \text{Nuts}_{\text{LB}} + 2 \leq \text{Cap}_{\text{LB}} .$$

- ♦ Suppose: there are 4 workers to install wheels, but they share one tool to put the axle in place.

→ We need a **disjunctive constraint** to say that Axle_F and Axle_B must not overlap in time; one must precede the other; so:

$$(\text{Axle}_F + 10 \leq \text{Axle}_B) \textbf{ or } (\text{Axle}_B + 10 \leq \text{Axle}_F)$$

Example: car assembly scheduling

- ♦ The *inspection* comes last and takes 3 minutes. →
For each variable except *Inspect*, we add a constraint of the form $X + d_x \leq \text{Inspect}$.
- ♦ Suppose there is a requirement to get the whole assembly done in 30 minutes.
 - Can be achieved by limiting the domain of all variables: $D_i = \{1, 2, 3, \dots, 27\}$.
- This problem is trivial to solve, but CSPs have been applied to job-shop scheduling problems like this with thousands of variables.
- In some cases, there are complicated constraints that are difficult to specify in the CSP formalism, and more advanced planning techniques are used.

Variations on the CSP Formalism

- Simplest kind of CSP: variables have **discrete, finite domains**. E.g.:
 - ♦ Map-colouring
 - ♦ Scheduling with time limits.
 - ♦ 8-Queens problem: each of the variables $Q_1, \dots, Q_8$ takes a value in $D_i = \{1, 2, 3, 4, 5, 6, 7, 8\}$
- A discrete domain can be **infinite**. E.g. if no deadline on the job-scheduling problem, there would be an infinite number of start times for each variable.
- For variables with infinite domains:
 - ♦ The constraints cannot be described by enumerating all combinations of possible values.
 - ♦ Need for a **constraint language**. E.g. to express a constraint like $T_1 + d_1 \leq T_2$ on variables T_1 and T_2 .

Variations on the CSP Formalism

- CSPs with infinite domains are also common:
 - ♦ E.g. the scheduling of experiments on the Hubble Space Telescope.
 - ♦ Field of Operations Research.
 - ♦ The best-known category of such CSPs is that of **linear programming** problems, where constraints must be linear equalities or inequalities.
- CSPs based on the types of constraints:
 - ♦ **Unary constraint:** simplest type; restricts the value of a constraint to a single variable. E.g. $\langle (SA), SA \neq \text{green} \rangle$.
 - ♦ **Binary constraint:** relates two variables. E.g. $SA \neq \text{NSW}$.
 - A binary CSP is one with only binary constraints and can be represented as a constraint graph.

Variations on the CSP Formalism

- A **ternary constraint** involves three variables, e.g. *between*(X, Y, Z).
- A **global constraint** involves any number of variables.
 - ♦ E.g. *Alldiff* CSP: all the variables involved in the constraint must have different values.
 - ♦ E.g. Sudoku: all variables in a row or column or box must satisfy an *Alldiff* constraint.
 - ♦ **Cryptarithmic** (*Alldiff*) puzzles: each letter in the puzzle represents a different digit. Eg. *Alldiff* (F, T, U, W, R, O).

$$\begin{array}{r} \text{T W O} \\ + \text{T W O} \\ \hline \text{= F O U R} \end{array}$$

Variations on the CSP Formalism

- The addition constraints for this puzzle can be written as:

$$O + O = R + 10 * C_1$$

$$C_1 + W + W = U + 10 * C_2$$

$$C_2 + T + T = O + 10 * C_3$$

$$C_3 = F$$

where C_1 , C_2 , and C_3 are auxiliary variables representing the digit (1 or 0) carried over into the 10s, 100s, or 1000s.


These constraints can be represented in a **constraint hypergraph**.

- A hypergraph consists of ordinary nodes (the circles in the figure) and hypernodes (the squares), which represent n-ary constraints.
- Note that **any CSP can be transformed into one with only binary constraints**. (Exercise)

Variations on the CSP Formalism

$$\begin{array}{r} T \ W \ O \\ + \ T \ W \ O \\ \hline F \ O \ U \ R \end{array}$$

(a)



(b)

Figure 6.2 (a) A cryptarithmic problem. Each letter stands for a distinct digit; the aim is to find a substitution of digits for letters such that the resulting sum is arithmetically correct, with the added restriction that no leading zeroes are allowed. (b) The constraint hypergraph for the cryptarithmic problem, showing the *Alldiff* constraint (square box at the top) as well as the column addition constraints (four square boxes in the middle). The variables *C*₁, *C*₂, and *C*₃ represent the carry digits for the three columns.

Variations on the CSP Formalism

- **Preference constraints** indicate which solutions are preferred. E.g., in a university class-scheduling problem:
 - ♦ There are **absolute constraints** that no professor can teach two classes at the same time.
 - ♦ **Preference constraints** may also be allowed : Prof. A prefers teaching in the morning, but Prof. B prefers teaching in the afternoon.
 - ➔ A schedule that has Prof. A teaches at 2 p.m. would be an allowable solution but not an optimal one.
- **Preference constraints can often be encoded as costs on individual variable assignments**
 - ♦ E.g., assigning an afternoon slot for Prof. A costs 2 points against the overall objective function, but a morning slot costs 1.
 - ♦ As such, **CSPs with preferences can be solved with optimization search methods**, either path-based or local.
 - ♦ We call such a problem a **constraint optimization problem (COP)**.

Outline

- Defining Constraint Satisfaction Problems
 - ♦ Example problem: Map colouring
 - ♦ Example problem: Job-shop scheduling
 - ♦ Variations on the CSP formalism
- Constraint Propagation: Inference in CSPs
 - ♦ Node consistency
 - ♦ Arc consistency
 - ♦ Path consistency
 - ♦ K -consistency
 - ♦ Global constraints
 - ♦ Sudoku example

Constraint Propagation: Inference in CSPs

- In CSPs there is a choice:
 - ♦ an algorithm can search (choose a new variable assignment from several possibilities)
 - ♦ and/or do a specific type of **inference** called **constraint propagation**: using the constraints to reduce the number of legal values for a variable, which in turn can reduce the legal values for another variable, and so on.
- The key idea is **local consistency** which is of different types: Node consistency, Arc consistency, Path consistency, *K*-consistency.

Node consistency

- A single variable (a node in the CSP network) is **node-consistent** if all the values in the variable's domain satisfy the variable's unary constraints.
 - ♦ In the variant of Australia map-colouring problem where South Australians dislike green, the variable SA starts with domain {red , green, blue}, and is made node-consistent by eliminating green → SA then has the reduced domain {red , blue}.
 - ♦ A network is node-consistent if every variable in it is node-consistent.
- It is always possible to eliminate all the unary constraints in a CSP by running node consistency.

Arc consistency

- A variable in a CSP is **arc-consistent** if every value in its domain satisfies the variable's binary constraints.
- More formally, X_i is arc-consistent with respect to variable X_j if $\forall d_i \in D_i, \exists d_j \in D_j$ s. t.

binary constraint on $\text{arc}(X_i, X_j)$ is satisfied

- ♦ E.g., consider the constraint $Y = X^2$ where the domain of X and Y is the set of digits $\{0, 1, 2, 3, 4, 5, 6, 7, 8, 9\}$.
 - This constraint can be written explicitly as:
 $\langle (X, Y), \{(0, 0), (1, 1), (2, 4), (3, 9)\} \rangle$
 - To make X arc-consistent with respect to Y , we reduce X 's domain to $\{0, 1, 2, 3\}$.
 - If we also make Y arc-consistent with respect to X , then Y 's domain becomes $\{0, 1, 4, 9\}$ and the whole CSP becomes arc-consistent

Arc consistency

- Note that arc consistency can do nothing for the Australia map-colouring problem.
 - ♦ E.g., consider the following inequality constraint on (SA,WA):
 $\{(\text{red}, \text{green}), (\text{red}, \text{blue}), (\text{green}, \text{red}), (\text{green}, \text{blue}), (\text{blue}, \text{red}), (\text{blue}, \text{green})\}.$
 - ♦ No matter what value you choose for SA (or for WA), there is a valid value for the other variable.
 - ♦ So applying arc consistency has no effect on the domains of either variable.

Arc consistency: AC-3

- AC-3: most popular algorithm for arc consistency.
- AC-3 maintains a queue of arcs to consider to make every variable arc-consistent. (Actually, the order of consideration is not important, so the data structure is really a set, but traditionally called a queue.)
- Initially, the queue contains all the arcs in the CSP.
- AC-3 then pops off an arbitrary arc (X_i, X_j) from the queue and makes X_i arc-consistent with respect to X_j .
- If this leaves D_i unchanged,
 - ♦ then the algorithm just moves on to the next arc.
- else, if it revises D_i (i.e. makes D_i smaller), so we add to the queue all arcs (X_k, X_i) where X_k is a neighbour of X_i . This is needed because the change in D_i might enable further reductions in the domains of D_k , even if we have previously considered X_k .

Arc consistency: AC-3

- If D_i is revised down to the empty set,
 - ♦ then we know the whole CSP has no consistent solution, and AC-3 returns failure.
 - ♦ else, we keep checking, trying to remove values from the domains of variables until no more arcs are in the queue.
 - At this point, we end up with a CSP that is equivalent to the original CSP—both have the same solutions.
 - The arc-consistent CSP will (in most cases) be faster to search because its variables have smaller domains.
- It is possible to extend the notion of arc consistency to handle n-ary (not just binary) constraints, called generalised arc consistency (or hyperarc consistency).

Arc consistency: AC-3

function AC-3(*csp*) **returns** false if an inconsistency is found and true otherwise

inputs: *csp*, a binary CSP with components (X , D , C)

local variables: *queue*, a queue of arcs, initially all the arcs in *csp*

while *queue* is not empty **do**

$(X_i, X_j) \leftarrow \text{REMOVE-FIRST}(\text{queue})$

if REVISE(*csp*, X_i , X_j) **then**

if size of $D_i = 0$ **then return** false

for each X_k **in** $X_i.\text{NEIGHBORS} - \{X_j\}$ **do**

 add (X_k, X_i) to *queue*

return true

function REVISE(*csp*, X_i , X_j) **returns** true iff we revise the domain of X_i

revised $\leftarrow$ false

for each x **in** D_i **do**

if no value y in D_j allows (x,y) to satisfy the constraint between X_i and X_j **then**

 delete x from D_i

revised $\leftarrow$ true

return *revised*

Figure 6.3 The arc-consistency algorithm AC-3. After applying AC-3, either every arc is arc-consistent, or some variable has an empty domain, indicating that the CSP cannot be solved. The name “AC-3” was used by the algorithm’s inventor (Mackworth, 1977) because it’s the third version developed in the paper.

Recall

- It is also possible to transform all n-ary constraints into binary ones.
- It is common to define CSP solvers that work with only binary constraints.
- Assumption made in the sequel.

Path consistency

- Arc consistency can go a long way toward reducing the domains of variables, sometimes every domain to one value or a domain as empty set.
- For other networks, arc consistency fails to make enough inferences. E.g. map-colouring problem on Australia.
- Arc consistency tightens down the domains (unary constraints) using the arcs (binary constraints).
- To make progress on problems like map colouring, we need a stronger notion of consistency.
- **Path consistency** tightens the binary constraints by using implicit constraints that are inferred by looking at triples of variables.

Path consistency

- A two-variable set $\{X_i, X_j\}$ is **path-consistent** with respect to a third variable X_m if, for every assignment $\{X_i = a, X_j = b\}$ consistent with the constraints on $\{X_i, X_j\}$, there is an assignment to X_m that satisfies the constraints on $\{X_i, X_m\}$ and $\{X_m, X_j\}$.
 - This is called **path consistency** because it can be seen as looking at a path from X_i to X_j with X_m in the middle.
 - E.g. colouring the Australia map with two colours:
 - ♦ Let us make the set $\{WA, SA\}$ path consistent with respect to NT.
 - ♦ Let us start by enumerating the consistent assignments to the set. There are only two:
 $\{WA = \text{red}, SA = \text{blue}\}$ and $\{WA = \text{blue}, SA = \text{red}\}$.
 - We can see that with both of these assignments NT can be neither red nor blue, so no valid assignments for $\{WA, SA\}$.
- ➔ There can be no solution to this problem.

K-consistency

- Stronger forms of propagation can be defined with the notion of **k-consistency**.
- *A CSP is k-consistent if, for any set of $k-1$ variables and for any consistent assignment to those variables, a consistent value can always be assigned to any k^{th} variable.*
- Special cases of k-consistency:
 - ♦ 1-consistency is the same as node consistency.
 - ♦ 2-consistency is the same as arc consistency.
 - ♦ 3-consistency is the same as path consistency.
- A CSP is **strongly k-consistent** if it is k-consistent and is also $(k-1)$ -consistent, $(k-2)$ -consistent, . . . all the way down to 1-consistent.

K-consistency

- Suppose we have a CSP with n nodes and it is strongly n -consistent.
- The problem can then be solved as follows:
 - ♦ First, choose a consistent value for X_1 .
 - ♦ Then one is guaranteed to be able to choose a value for X_2 because the graph is 2-consistent,
 - ♦ .. And for X_3 because it is 3-consistent, and so on.
 - ♦ For each variable X_i , we need only search through the d values in its domain D_i to find a value consistent with $X_1, \dots, X_{i-1}$.
- We are guaranteed to find a solution in time $O(n^2 d)$ (where $d = \max_i |d_i|$)
- However, any algorithm for establishing n -consistency must take time exponential in n in the worst case. Same for space complexity.
- In practice, practitioners commonly compute 2-consistency and less commonly 3-consistency.

Global constraints

- Remember: a **global constraint** is one involving an arbitrary number of variables (not necessarily all variables).
- These constraints can be handled by special-purpose algorithms that are more efficient than the general-purpose methods described so far.
 - ♦ E.g. ***Alldiff* problems**: all variables must have distinct values
 - ♦ A simple detection of inconsistency for *Alldiff* constraints: if it is an m -variable constraint, and there are n possible distinct values altogether, and $m > n$, then the constraint cannot be satisfied.
 - ♦ ➔ simple algorithm:
 1. Remove any variable X with singleton value v
 2. Remove v from the domains of the remaining variables
 3. Keep repeating steps 1 and 2; if empty domain reached or number of variables > number of available values, then **inconsistency detected**.
 - ♦ E.g. {WA=red , NSW=red} and SA, NT, Q to be coloured!

Example problem: Map colouring




Figure 6.1 (a) The principal states and territories of Australia. Coloring this map can be viewed as a constraint satisfaction problem (CSP). The goal is to assign colors to each region so that no neighboring regions have the same color. (b) The map-coloring problem represented as a constraint graph.

Global constraints

- Other global constraint: **resource constraint**, also called the *atmost constraint*.
 - ♦ E.g., in a scheduling problem, let $P_1, \dots, P_4$ denote the numbers of workers assigned to each of four tasks.
 - ♦ Constraint: no more than 10 workers are assigned in total is written as $Atmost(10, P_1, P_2, P_3, P_4)$.
 - ♦ One can detect inconsistency: check the sum of the minimum values of the current domains of the variables.
 - E.g. if variables domain is $\{3, 4, 5, 6\}$, then *Atmost* constraint cannot be satisfied.
 - ♦ Can also enforce consistency: delete the max value of any domain if it is not consistent with the min values of the other domains.
 - E.g. if variables domain is $\{2, 3, 4, 5, 6\}$, then the values 5 and 6 can be deleted from each domain.

Global constraints

- For large resource-limited problems with integer values:
 - ♦ E.g. logistics problems of moving thousands of people in hundreds of vehicles
- Not possible to represent the domain of a variable as a large set of integers and gradually reduce it by consistency-checking methods.
- Domains are represented by upper and lower bounds and **bounds propagation** is performed.
 - ♦ E.g., airline-scheduling problem: suppose for two flights F1 and F2, the planes have capacities 165 and 385, resp.
 - ♦ The initial domains for the numbers of passengers on each flight are then $D1 = [0, 165]$ and $D2 = [0, 385]$.
 - ♦ Suppose the additional constraint: the two flights together must carry 420 people: $F1 + F2 = 420$.
 - ♦ Propagating bounds constraints, we reduce the domains to $D1 = [35, 165]$ and $D2 = [255, 385]$.

Sudoku example

	1	2	3	4	5	6	7	8	9
A			3		2		6		
B	9			3		5			1
C			1	8		6	4		
D			8	1		2	9		
E	7								8
F			6	7		8	2		
G			2	6		9	5		
H	8			2		3			9
I			5		1		3		

(a)

	1	2	3	4	5	6	7	8	9
A	4	8	3	9	2	1	6	5	7
B	9	6	7	3	4	5	8	2	1
C	2	5	1	8	7	6	4	9	3
D	5	4	8	1	3	2	9	7	6
E	7	2	9	5	6	4	1	3	8
F	1	3	6	7	9	8	2	4	5
G	3	7	2	6	8	9	5	1	4
H	8	1	4	2	5	3	7	6	9
I	6	9	5	4	1	7	3	8	2

(b)

Figure 6.4 (a) A Sudoku puzzle and (b) its solution.

Sudoku example

- Sudoku board: 81 squares, some of which are initially filled with digits from 1 to 9.
- Puzzle: fill in all the remaining squares such that no digit appears twice in any row, column, or 3×3 box.
- A row, column, or box is called a **unit**.
- ➔ can be seen as a CSP with 81 variables, one for each square.
- Variable names A1 through A9 for the top row (left to right), down to I1 through I9 for the bottom row.
- Empty squares have the domain $\{1, 2, 3, 4, 5, 6, 7, 8, 9\}$
- Prefilled squares have a domain consisting of a single value.

Sudoku example

- In addition, there are 27 different Alldiff constraints: one for each row, column, and box of 9 squares.
 - ♦ *Alldiff (A1,A2,A3,A4,A5,A6, A7, A8, A9)*
 - ♦ *Alldiff (B1,B2,B3,B4,B5,B6,B7,B8,B9)*
 - ♦ . . .
 - ♦ *Alldiff (A1,B1,C1,D1,E1, F1,G1,H1, I1)*
 - ♦ *Alldiff (A2,B2,C2,D2,E2, F2,G2,H2, I2)*
 - ♦ . . .
 - ♦ *Alldiff (A1,A2,A3,B1,B2,B3,C1,C2,C3)*
 - ♦ *Alldiff (A4,A5,A6,B4,B5,B6,C4,C5,C6)*
 - ♦ . . .

Arc Consistency on Sudoku

- Applying arc consistency:
 - ♦ Assume that the *Alldiff* constraints have been expanded into binary constraints (e.g. $A1 \neq A2$)
 - ♦ So we can apply the AC-3 algorithm directly.
 - ♦ In the example, consider variable E_6
 - From constraints in the box: we can remove 1, 2, 7, 8 from E_6 's domain.
 - From constraints in its column: remove 5, 6, 2, 8, 9, and 3. $\rightarrow E_6$ has singleton domain $\{4\}$. E_6 now known.
 - ♦ Then consider variable I_6
 - Applying arc consistency in its column, remove 5, 6, 2, 4, 8, 9, and 3.
 - By arc consistency on row, remove 1 $\rightarrow$ domain of I_6 is $\{7\}$.
 - Now we can infer that $A6$ must be 1.

Arc Consistency on Sudoku

- AC-3 works only for the easiest Sudoku puzzles.
- Slightly harder ones can be solved by PC-2 (for it achieves path consistency), but at a greater computational cost: there are 255,960 different path constraints to consider in a Sudoku puzzle.
- To solve the hardest puzzles, need to apply more complex inference strategies.
- The “**naked triples**” strategy works as follows:
 - ♦ In any unit (row, column or box), find three squares that each have a domain that contains the same three numbers or a subset thereof and **no other variables in that scope** can take those values. E.g. {1, 8}, {3, 8}, and {1, 3, 8}.
 - ♦ Here we don't know which square contains 1, 3, or 8.
 - ♦ BUT we know we can remove 1, 3, and 8 from the domains of every *other* square in the unit.

Sudoku and CSP

- Note that all that has been done is not specific to Sudoku (including the “naked triples” strategy)
- We do have to say that:
 - ♦ there are 81 variables,
 - ♦ their domains are the digits 1 to 9, and
 - ♦ there are 27 Alldiff constraints.
- Beyond that, all the strategies—arc consistency, path consistency, etc.—apply generally to all CSPs, not just to Sudoku problems.
- The power of the CSP formalism: for each new problem area, we only need to define the problem in terms of constraints; then the general constraint-solving mechanisms can take over.


Outline

- Defining Constraint Satisfaction Problems
 - ♦ Example problem: Map colouring
 - ♦ Example problem: Job-shop scheduling
 - ♦ Variations on the CSP formalism
- Constraint Propagation: Inference in CSPs
 - ♦ Node consistency
 - ♦ Arc consistency
 - ♦ Path consistency
 - ♦ K-consistency
 - ♦ Global constraints
 - ♦ Sudoku example

Outline

- Backtracking Search for CSPs
 - ◆ Variable and value ordering
 - ◆ Interleaving search and inference
 - ◆ Intelligent backtracking: Looking backward
- Local Search for CSPs
- The Structure of Problems

Slides based on the textbook



- Russel, S. and Norvig, P. (2020) Artificial Intelligence, A Modern Approach (4th Edition), Pearson Education Limited.