

AI224: Operating Systems — Lecture 3

Dr. Karim Lounis

Jan - May 2025



Process Management

Processes

Definition

Process. Is an instance of a computer program that is being executed.

- A process consists of:

- **An address space.** Memory locations where the program's code, variables, and function frames are stored.
- **The CPU state.** The values of CPU registers (e.g., PC, SP, etc) during the execution of the process (and while it is waiting/blocked).
- **A set of OS resources.** Represents the logical and physical resources the process holds on, e.g., open files, network connections, ...
- **A lifetime.** Since its creation, a process goes through different phases: created-executed-[interrupted-resumed]-terminated.

A program is a passive entity, whereas a process is its active form.

During its lifetime, a process is identified by a unique [temporary] number called PID (Process Identifier) (e.g., in GNU/Linux [ps](#) [-aux](#) or [top](#)).

Processes (top)

Output of \$ top command run on MacOS Monterey (Apple M1):

Processes: 492 total, 2 running, 498 sleeping, 2315 threads
 Load Avg: 1.58, 1.48, 1.57 CPU usage: 1.64k user, 3.88k sys, 94.47% idle SharedLibs: 328M resident, 75M data, 28M linkedit. MemRegions: 32580 total, 12802M resident, 189M private, 1823M shared.
 PhysMem: 7340M used (1287M wired), 799M unused. VM: 2317 vszize, 3823M framework vszize, 31894(0) swpains, 33552(0) swappouts. Networks: packets: 236689/234M in, 182487/34M out.
 Dicks: 648222/270 read, 228893/4050M written.

#	COMMAND	%CPU	TIME	MEM	VSZ	SHR	WIR	RES	MEM	PURG	CHMRS	PGPR	PPID	STATE	BOOSTS	%CPU	MEM	NCPU_OTHERS	UID	FAULTS	COM	MSGSENT	MSGRECV	SYSDSD	SYSMACH	CSW	PAGEIN	IDLEW
0	kernel_task	15.6	04:01:15	583/8	0	0	18M	08	08	0	0	0	running	0(0)	0.00000	0.00000	0	29394	0	5874943	3539373	0	0	18437867	0	1835923		
375	WindowServer	0.1	04:03:30	21	5	1990	397M+	13M+	08	61M	375	1	sleeping	*0(1)	0.06698	0.33386	88	895085+	8651	2666252+	181269+	5285241+	9465263+	2836761+	2170	88578+		
3772	top	5.4	00:01:02	1/1	0	28	5937K	08	08	3772	632	running	*0(1)	0.00000	0.00000	0	12968+	78	981329+	498659+	58418+	586515+	318+	15	1	0		
394	nsurlsession	3.4	00:00:15	7	6	109	5857K	08	08	1472K	394	1	sleeping	*0(3)	0.00000	0.78562	242	17676+	132	7974+	4418+	495898+	17668+	118953+	289	4644+		
732	Google Chrome	1.7	00:51:06	12	1	146	43M+	08	08	15M	563	563	sleeping	*0(3)	0.00000	0.00000	581	43445+	648	498875+	367936+	1487859+	855655+	399259+	265	18034+		
833	distnoted	1.0	00:53:43	2	1	482	2417K	08	08	336K	533	1	sleeping	*0(1)	0.00000	0.95354	581	745	46	218707+	42674+	335965+	187221+	45564+	1	0		
3715	com.apple.St	0.9	00:02:27	2	1	35	5569K	08	08	3715	1	sleeping	*0(0)	0.00000	0.83908	242	7588+	68	4487+	3897+	14835+	9331+	3861+	11	0			
494	Terminal	0.8	00:01:70	9	4	274	92M	21M	08	6288K	594	1	sleeping	*0(104)	0.05470	0.00000	581	15926	385	16188+	5184+	26437+	32628+	18855+	987	199+		
563	Google Chrome	0.8	02:05:26	39	1	929	274M	08	08	115M	563	1	sleeping	*0(398)	0.00000	0.00000	581	325224+	2455	216891+	1142557+	239489+	7213128+	1339609+	4198	1669+		
593	AdobeReader	0.8	00:31:37	17	4	233	183M	08	08	73M	593	1	sleeping	*0(153)	1.74444	0.00000	581	15525	745	37218+	5879	284898+	3643855+	489734+	266	119801+		
355	distnoted	0.5	00:16:49	2	1	194	1585K	08	08	416K	355	1	sleeping	*0(1)	0.00000	0.49056	241	593	46	223959+	33842+	190839	291596+	22687+	6	0		
359	AdskLicensin	0.4	00:15:96	12	0	47	12M	08	08	5658K	359	1	sleeping	*0(1)	0.00000	0.00000	499	7887	69	4468	1335	975288+	4344	831828+	1839	223846+		
491	mobileassetd	0.3	00:19:55	2	4	158	11M	08	08	4864K	491	1	sleeping	*0(39)	0.20788	0.25668	0	17615	637	2357+	2997+	93189+	8562+	29789+	111	252		
1558	com.apple.Dr	0.2	00:00:12	3	2	5186	6177K	08	08	1872K	1558	1	sleeping	*0(1)	0.00000	0.00000	278	488	32	43831+	33138+	67521+	22899+	4636+	51	3344+		
3773	screenapp	0.2	00:00:45	4	2	145	8434K	752K	08	681	681	1	sleeping	*0(329)	0.00000	0.00000	581	6872+	773+	3671+	1863+	2489+	8698+	2808+	65	2		
1085	Google Chrome	0.2	00:18:54	19	2	282+	96M+	08	08	33M	563	563	sleeping	*0(6)	0.00000	0.00000	581	59968+	691	57111+	27864+	175374+	164358+	181587+	0	3672+		
1881	com.apple.Am	0.2	00:01:22	4	2	72	3585K	08	08	1792K	1881	1	sleeping	*0(11)	0.13531	0.00000	0	877	89	6781+	3136+	21681+	1581+	11978+	67	3		
829	rdarCF Helpe	0.2	00:10:14	18	2	154	43M	08	08	288K	768	768	sleeping	*0(15)	0.00000	0.00000	581	7431	744	2146	227193+	280648+	538644+	335892+	374	41612+		
389	fsventnoted	0.1	00:00:51	18	1	273	4545K	08	08	736K	389	1	sleeping	*0(1)	0.00000	0.00000	0	38816+	62	31144+	12861+	37421+	54778+	61551+	18	5783+		
368	corebrightness	0.1	00:05:75	4	3	138	4545K	08	08	2948K	368	1	sleeping	*0(1)	0.14681	0.00000	0	13367+	97	19514+	26739+	117145+	42183+	33886+	41	1869		
380	logd	0.1	00:16:26	4	3	1635	28M	08	08	17M	305	1	sleeping	*0(1)	0.00000	0.01535	0	62827	69	33618+	38741+	292117+	49485+	113687+	46184	143		
752	Google Chrome	0.1	00:06:08	13	1	122	25M+	08	08	18M	563	563	sleeping	*0(3)	0.00000	0.00000	581	13378+	661	184867+	56497+	125958+	237759+	92682+	8	1929+		
508	softwarep	0.1	00:28:07	5	3	158	28M	08	08	2448K	508	1	sleeping	*0(32)	0.00000	0.15176	288	16888	168	152848+	88788+	168668+	122827+	45766+	1139	179		
316	powerd	0.1	00:04:00	4	3	136	4865K	08	08	992K	316	1	sleeping	*0(1)	0.00000	0.06698	0	24402+	186	58981+	28699+	85481+	96755+	35748+	119	188		
752	Google Chrome	0.1	01:06:08	22	5	347	485M	2048K	08	56M	563	563	sleeping	*0(15)	0.00000	0.00000	581	151513	730	856575	867438+	849988+	2987889+	783266+	657	11812		
1522	PerfPowerSer	0.0	00:07:17	5	3	885	18M	256K	08	6288K	1562	1	sleeping	*0(29)	0.00000	0.00000	0	6766	156	454083	234684+	57898+	290808	68896+	182	44		
682	Finder	0.0	00:28:84	2	2	543+	81M+	08	08	33M	682	1	sleeping	*0(282)	0.00000	0.00000	581	189015+	465	169982	88984+	424882+	438646+	628517+	277238+	246	4278	
481	TouchBarServ	0.0	00:17:24	4	2	289	17M	784K	08	6484K	481	1	sleeping	*0(1)	0.00000	0.08297	0	5857	167	121337+	78233+	438646+	628517+	277238+	246	4278		
771	Adobe Crash	0.0	00:02:70	4	3	54	4897K	08	08	2176K	771	593	sleeping	*0(2)	0.02827	0.00000	581	4998+	68	24861+	4879+	37187+	64357+	57898+	0	2491+		
733	Google Chrome	0.0	00:01:04	9	1	110	21M	08	08	16M	563	563	sleeping	*0(3)	0.00000	0.00000	581	11458	646	8879	5135	1435+	18548+	9962+	21	733+		
1219	Google Chrome	0.0	00:18:50	16	1	271+	104M+	08	08	63M	563	563	sleeping	*0(6)	0.00000	0.00000	581	58085+	695	51668+	32866	136927+	143223+	67154+	44	748+		
395	distnoted	0.0	00:02:05	2	1	32	1185K	08	08	354K	395	1	sleeping	*0(1)	0.00000	0.03326	0	488	46	88	16281+	97785+	32772+	17387+	0	0		
479	distnoted	0.0	00:01:03	2	1	32	1121K	08	08	448K	479	1	sleeping	*0(1)	0.00000	0.02321	205	498	46	92	16279+	97475+	32772+	17343+	0	0		
495	distnoted	0.0	00:01:43	2	1	29	1137K	08	08	464K	495	1	sleeping	*0(1)	0.00000	0.02291	292	460	46	89	16277+	81583+	32798+	17276+	0	0		
323	AdskAccessSe	0.0	00:02:34	6	1	77	8878K	08	08	6820K	323	1	sleeping	*0(1)	0.00000	0.00000	0	18472	116	9528+	4734+	48807+	8118+	41285+	2996	18676+		
466	distnoted	0.0	00:01:43	2	1	29	1185K	08	08	448K	466	1	sleeping	*0(1)	0.00000	0.02156	88	740	46	90	16276+	81578+	32798+	17295+	0	0		
338	nds	0.0	00:16:49	8	5	322	37M	08	08	8268K	338	1	sleeping	*0(1)	0.00000	0.00000	0	58358	143	24430	18965	518974+	74993	19111+	284	4396+		

Here the first process is **kernel-task** with PID=0. Also, the process related to the **top** command, which displayed this list of processes, has PID=3772.

Processes (top)

Output of `$ top` command run on GNU/Linux Debian:

```
top - 06:31:08 up 1 min, 1 user, load average: 0.30, 0.11, 0.04
Tasks: 229 total, 1 running, 228 sleeping, 0 stopped, 0 zombie
%Cpu(s): 0.0 us, 0.0 sy, 0.0 ni, 99.9 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
MiB Mem : 3925.8 total, 2333.6 free, 783.2 used, 809.1 buff/cache
MiB Swap: 976.0 total, 976.0 free, 0.0 used. 2878.2 avail Mem
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
1	root	20	0	0	0	0	I	0.3	0.0	0:00.27	kworker/0:1-events
1301	jeffery	20	0	4587260	197060	107568	S	0.3	4.9	0:04.56	gnome-shell
1794	jeffery	20	0	402080	49320	39200	S	0.3	1.2	0:00.75	gnome-terminal-
2118	jeffery	20	0	10100	3352	2616	R	0.3	0.1	0:00.15	top
1	root	20	0	164824	9956	7260	S	0.0	0.2	0:01.98	systemd
2	root	20	0	0	0	0	S	0.0	0.0	0:00.03	kthreadd
3	root	0	-20	0	0	0	I	0.0	0.0	0:00.00	rcu_gp
4	root	0	-20	0	0	0	I	0.0	0.0	0:00.00	rcu_par_gp
5	root	20	0	0	0	0	I	0.0	0.0	0:00.00	kworker/0:0-events
6	root	0	-20	0	0	0	I	0.0	0.0	0:00.00	kworker/0:0H-events_highpri
7	root	20	0	0	0	0	I	0.0	0.0	0:00.01	kworker/u16:0-flush-254:0
8	root	0	-20	0	0	0	I	0.0	0.0	0:00.00	mm_percpu_wq
9	root	20	0	0	0	0	S	0.0	0.0	0:00.00	rcu_tasks_rude
10	root	20	0	0	0	0	S	0.0	0.0	0:00.00	rcu_tasks_trace
11	root	20	0	0	0	0	S	0.0	0.0	0:00.01	ksoftirqd/0
12	root	20	0	0	0	0	I	0.0	0.0	0:00.25	rcu_sched
13	root	rt	0	0	0	0	S	0.0	0.0	0:00.03	migration/0
15	root	20	0	0	0	0	S	0.0	0.0	0:00.00	cpuhp/0
16	root	20	0	0	0	0	S	0.0	0.0	0:00.00	cpuhp/1
17	root	rt	0	0	0	0	S	0.0	0.0	0:00.03	migration/1
18	root	20	0	0	0	0	S	0.0	0.0	0:00.01	ksoftirqd/1
19	root	20	0	0	0	0	I	0.0	0.0	0:00.00	kworker/1:0-events
20	root	0	-20	0	0	0	I	0.0	0.0	0:00.00	kworker/1:0H-events_highpri

Here the first process is `systemd` (PID=1), which used to be `init` (process responsible for starting all other processes). Also, the process related to the `top` command, which displayed this list of processes, has PID=2118.

Types of Processes

Two types of processes: system processes that execute system codes, and user processes that execute user codes (but treated in the same way).

System processes are executed in kernel mode (master mode — system mode) whereas user processes are executed in user mode (slave mode).

E.g., [Google Chrome](#), [Texteditor](#), [IDE](#), [etc](#) processes are user processes, whereas [system update](#), [demons](#), [event logger](#), [etc](#) are system processes.

Various terms are used (interchangably): process, job, task (batch mode).

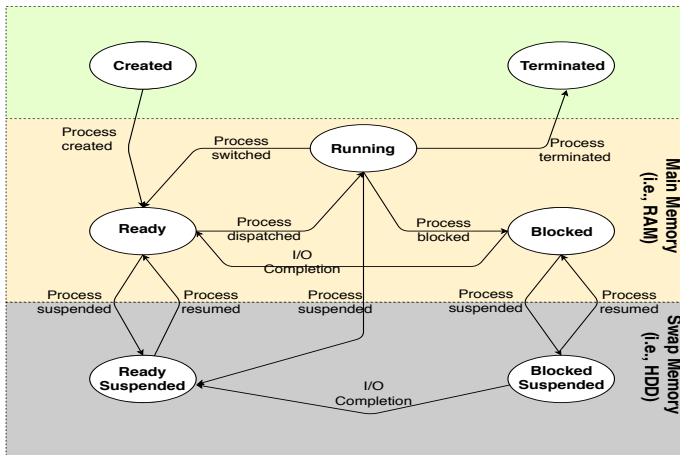
- **Process.** Defined ealier (cf., previous slide).
- **Task.** Could be a process or a thread performing some computation (e.g., printing, downloading, updating, etc). When multiple tasks are executed, that is multitasking — multithreading or multiprocessing.

[In Java, when a thread is running, it is called a task.](#)

- **Job.** A complete unit of work. It could be a set of tasks.

Process States (Process State Graph)

During its lifetime, a process transits from one state to another as follows:



Process States (Process State Graph)

The different states are defined as follows:

- **New.** Just been created.
- **Ready.** Having all needed resources and waiting (ready queue) for CPU to be allocated.
- **Running.** Being executed on the CPU.
- **Blocked.** Waiting (blocked queue) for another process to terminate (e.g., I/O routines) or an event to happen (e.g., signal reception).
- **Suspended.** Swapped out of RAM for some reason, e.g.,:
 - Swapped out of RAM to free some space for other process.
 - Intentionally suspended by superuser ([8am-6pm]) to save CPU cycles.
 - Swapped out of RAM after being longly waiting for an event to happen.
 - Swapped out by the parent process.
- **Terminated.** The process has finished its task and/or have been killed (e.g., in Unix-like operating systems the command `kill -9 2730` is used to terminate the process which PID is 2730).

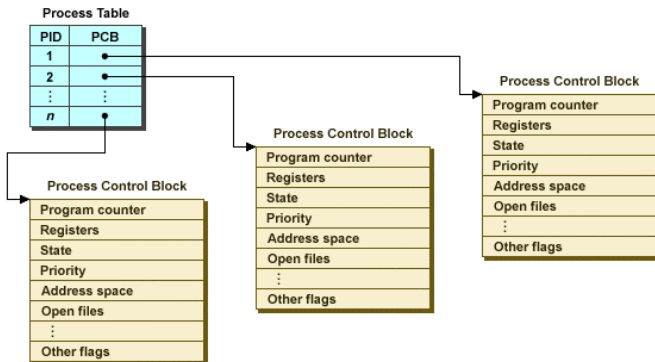
Process States (Process State Graph)

On Unix-like operating systems, processes could be in:

- **Running** — **R**. Currently executing on the CPU.
- **Sleeping** — **S**. Waiting for a resource to be available.
- **Stopped** — **T**. Suspended and can be resumed.
- **Idle** — **I**. Idle kernel threads — waiting for a task to be assigned.
- **Zombie** — **Z**. Processes finished, but father not.
- **Dead** — **X**. Process has terminated (not visible in the list).

Process Control Block

Every process is represented by a data structure called PCB (Process Control Block) that stores information about the process last state:



Enough information to resume execution and track process.

Process Control Block (GNU/Linux ps)

```

jeffery@pc-debian: ~
jeffery@pc-debian:~$ cat dummy.c
void main()
{
    while(1==1);
    return;
}
jeffery@pc-debian:~$ gcc dummy.c -o dummy.exe
jeffery@pc-debian:~$ ls -l dummy.c dummy.exe
-rw-r--r-- 1 jeffery jeffery 39 Feb 5 08:16 dummy.c
-rwxr-xr-x 1 jeffery jeffery 9288 Feb 5 08:21 dummy.exe
jeffery@pc-debian:~$ █

jeffery@pc-debian:~$ ./dummy.exe &
[1] 2427
jeffery@pc-debian:~$ ps
  PID TTY          TIME CMD
 2267 pts/3        00:00:00 bash
 2427 pts/3        00:00:02 dummy.exe
 2429 pts/3        00:00:00 ps
jeffery@pc-debian:~$ kill -STOP 2427

[1]+  Stopped                  ./dummy.exe
jeffery@pc-debian:~$ kill -CONT 2427
jeffery@pc-debian:~$ kill -9 2427
jeffery@pc-debian:~$ █

jeffery@pc-debian: ~
jeffery@pc-debian:~$ ps -o pid,ppid,comm,args,stat,user,pri,policy,nlwp,etime,rss,vsz,sz -p 2427
  PID  PPID COMMAND          COMMAND          STAT USER  PRI POL NLWP  ELAPSED  RSS   VSZ   SZ
 2427  2267 dummy.exe          ./dummy.exe      R   jeffery  19 TS  1      07:59   404  1796  449
jeffery@pc-debian:~$ ps -o pid,ppid,comm,args,stat,user,pri,policy,nlwp,etime,rss,vsz,sz -p 2427
  PID  PPID COMMAND          COMMAND          STAT USER  PRI POL NLWP  ELAPSED  RSS   VSZ   SZ
 2427  2267 dummy.exe          ./dummy.exe      T   jeffery  19 TS  1      09:24   404  1796  449
jeffery@pc-debian:~$ ps -o pid,ppid,comm,args,stat,user,pri,policy,nlwp,etime,rss,vsz,sz -p 2427
  PID  PPID COMMAND          COMMAND          STAT USER  PRI POL NLWP  ELAPSED  RSS   VSZ   SZ
 2427  2267 dummy.exe          ./dummy.exe      R   jeffery  19 TS  1      09:59   404  1796  449
jeffery@pc-debian:~$ ps -o pid,ppid,comm,args,stat,user,pri,policy,nlwp,etime,rss,vsz,sz -p 2427
  PID  PPID COMMAND          COMMAND          STAT USER  PRI POL NLWP  ELAPSED  RSS   VSZ   SZ
jeffery@pc-debian:~$ █

```

Here, I create a program that spins around forever (dummy.c). Then, I run it, suspend it, resume it, kill it, and observe certain of its PCB's fields.

Process Creation and Termination (Unix-like Systems)

Creation. A process can create other processes (children).

- In Unix-like systems, a process executes the `fork()` system call.
The OS duplicates the parent PCB and makes appropriate changes.
- The parent process can continue its execution or wait (call `wait()`).
- A process can pass data to another process (using `pipes`).
- The address space can be, either a duplication of the parent space (same data and code), or a newly loaded program (call `execvp()`).
- Resources are inherited (e.g., open files, variables, privileges, ...).

Termination. Occurs as follows:

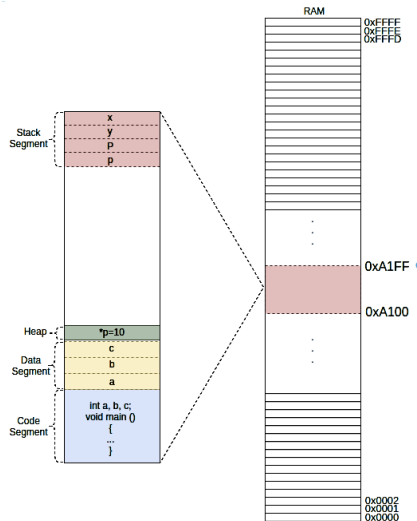
- A process executes the `exit()` system call.
- Running `exit()` returns an integer to the parent.
- A parent process can terminate a child using `kill()` system call.

Process Creation and Termination (Unix-like Systems)

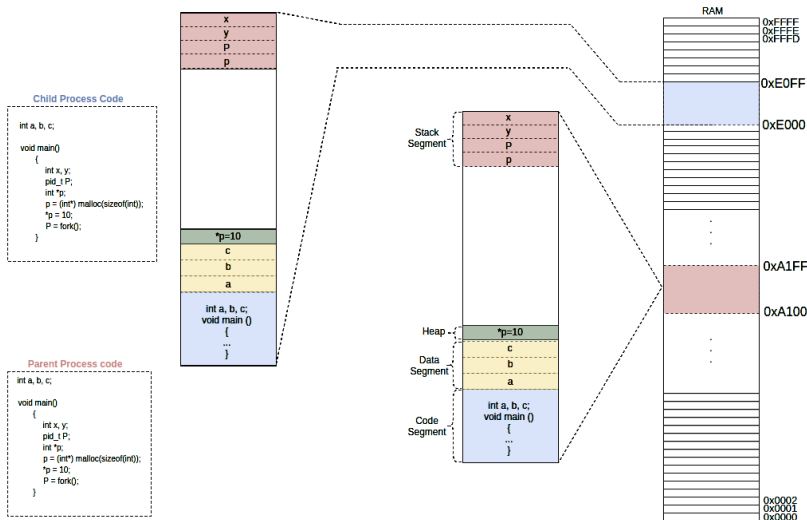
Parent Process code

```
int a, b, c;

void main()
{
    int x, y;
    pid_t P;
    int *p;
    p = (int*) malloc(sizeof(int));
    *p = 10;
    P = fork();
}
```



Process Creation and Termination (Unix-like Systems)



Process Creation and Termination (Unix-like Systems)

Exercise

Q1. Given the following parent process code snippet, how many processes would be created in total?

```
...  
for(unsigned int i=1; i<3; i++)  
{  
    fork();  
    printf("Welcome to ENSIA");  
}  
...
```

Q2. How many times the message “Welcome to ENSIA” gets printed?

Try it at: <https://www.programiz.com/c-programming/online-compiler/>

Concurrent Processes

Processes are said to be concurrent when:

- Multiple programs are loaded into the main memory, i.e., RAM.
- Their execution is happening in parallel or in fake-parallel.
- Their time durations or executions are overlapping or interleaving.

We can use PPG (Process Precedence Graph) to visualize execution order of concurrent processes. An execution can be sequential (serial) or parallel:




Figure: Processes Precedence Graph (Left) and its pseudocode (Right)

Interprocess Communication (a glance)

Processes can cooperate with each other for several reasons:

- **Information Sharing.** Collaborating on the same file.
- **Computational Speedup.** Parallel execution of different parts of a program (multiprocessors systems only).
- **Modularity.** An operating system can be constructed in a modular fashion. E.g., modern operating systems.
- **Convenience.** An individual users may run different processes.

Operating systems provides two fundamental mechanisms:

- 1 **Shared memory** Need a shared memory region (localhost):
 - Files: Three processes concurrently read/write into a same file (Lab 4)
 - Signals: 2 processes cooperating for a sum using wait/exit (Lab 4)
 - Pipes: Writing into a file and reading from the other end (Lab 5)
 - Variables: Multithreading program to crack password (Lab 6)
- 2 **Message passing.** Need a physical/logical communication link:
 - Sockets: Processes communicating while running on different PCs.
 - RPCs: Processes invoking functions implemented on different PCs.

Interrupts

Interrupts & Interrupt handlers

Interrupts

Definition

Interrupt. Is an event that alters the sequence of instructions executed by the CPU. A dedicated program called ISR (Interrupt Service Routine) or interrupt handler is executed as a consequence.

Each interrupt has its own interrupt handler (written in assembly language).

Definition

Interrupt handler. A.k.a., ISR (Interrupt Service Routine), is a low-level program that is executed in kernel mode to service the cause of the interrupt.

Interrupt handlers are located by an IVT (Interrupt Vector Table)¹ stored in the memory. [E.g., in x86 architectures, it is generally stored at address 0x0000-0x03FF.](#) The number of interrupt types is limited by the architecture.

¹Also check IDT (Interrupt Descriptor Table) used on modern [x86 architectures](#)

Types of Interrupts

Software-based Interrupts. When the system is interrupted following the execution of an instruction in a program code, the interrupt is said to be software-based. There are two types:

- 1 **Exceptions.** An exception happens when the current instruction perform an illegal action such as: division by 0, arithmetic overflow, buffer overflow, page fault, access protected memory space, or attempt to execute a privileged instruction (**process is terminated**).
- 2 **System calls.** A.k.a., traps. User program are executed in user mode. In this mode, programs do not have direct access to peripheral. They have to issue a **system call** which causes an interrupt (**INT instruction**).

Hardware-based Interrupts: These asynchronous interrupts are triggered by a hardware unit such as the timer/clock, keyboard, mouse, power button, ..., device controllers or DMA, to get the attention from the CPU. They can be maskable (MI) or non-maskable (NMI).

Types of Interrupts

- **Maskable.** These are interrupts that the CPU can choose to ignore based on its current state and priority. They are typically used for non-critical events that can be handled later without causing significant issues. E.g., Keyboard input, network packets arriving, etc.
- **Non-Maskable.** These are interrupts that the CPU cannot ignore. They demand immediate attention, often for critical events that could jeopardize system stability or integrity. E.g., hardware errors, system watchdog timer expiring, etc..

Types of Interrupts

Other types of interrupts²:

- **IPI (Inter-Processor Interrupt)**. Used between processors to communicate in a multiprocessor system, e.g., one processor core requests another processor core to stop when the system is shutting down by the first processor (\$utdown now).
- **SI (Spurious Interrupt)**. A.k.a., phantom or ghost interrupt, is an unwanted hardware interrupt, e.g., when an interrupt that has been signaled to the processor and it is no longer required (or interrupt source disappeared), or buggy hardware.

Interrupt Storm. Occurs whenever the operating system receives a large number of interrupts (from hardware) that consumes the majority of the processor's time spurious signals, interrupt handler execution took too long, ..., faulty drivers. The system becomes non-responsive.

²Also, you may wanna check raster interrupts.

System Calls

Definition

System Calls: Constitutes an interface between the user (programs) and the services that are made available by an operating system.

- A systems call is a programmatic way in which a computer program requests a service from the operating system.
- Each operating system has its own system calls (e.g., GNU/Linux has 300, FreeBSD has 500, and Windows 7 has 700).
- System calls point to specific programs called routines that are written in C, C++, or assembly language.
- Operating systems provide an API (Application Programming Interface) as an interface between user programs and system calls.

E.g., `printf()` or `cout` are APIs for calling functions from C libraries. These functions uses the `wirte()` system call to send data to the output.

System Calls

System calls can be grouped into six categories:

- 1 **Process Control:** End, abort, load, execute, create process, terminate process, get/set process attributes, wait for a time, wait for an event, signal an event, allocate and free memory.

e.g., *exec()*, *fork()*, *wait()*, *exit()*, *getuid()*, *getgid()*, *getpid()*, *getppid()*, *signal()*, *kill()*, *acquire_lock()*, *release_lock()*.

- 2 **File Management:** Create, delete, open, and close files, read, write on files, get/set files attributes.

e.g., *create()*, *open()*, *close()*, *read()*, *write()*, *lseek()*, *dup()*, *link()*, *unlink()*, *stat()*, *fstat()*, *access()*, *chmod()*, *chown*, *umask*, *ioctl()*.

System Calls

- ③ **Device Manipulation:** Request device, release device, read, write on device, get/set devices attributes.

e.g., *request()*, *release()*, *ioctl()*, *read()*, *write()*, *open()*, *close()*.

- ④ **Information Maintenance:** get/set time or data, get/set system data, get/set resource attributes.

e.g., *time()*, *date()*, *alarm()*.

- ⑤ **Communication:** Create, delete a network connection, send and receive messages, transfer status information.

e.g., *pipe()*, *mmap()*, *shm_open()*, *gethostid()*.

- ⑥ **Protection:** Set ownership and access rights on a resource.

e.g., *chmod()*, *umask()*, *chown()*.


Hardware Interrupt mechanism

How Hardware events (keyboard presses, incoming network packets, mouse movement, ...) are escalated to running programs?


There exists two ways:

- Polling: When OS periodically queries each device to see if new information is available (simple but high latency and CPU cycle wasting).
- Interrupts: Device sends a signal to **interrupt controller**³ to request attention, CPU preempts running process to handle device request.

After each (user mode) instruction is executed, the CPU checks whether the interrupt controller has any interrupt pending. If an interrupt is there, the current process is “interrupted”, and the appropriate handler is executed.

³PIC & APIC, generally, integrated circuits within the southbridge. 

Hardware Interrupts and Interrupt Service Routines



Hardware interrupts.

● End