

Data Mining

Data: Part 2

Mohammed Brahimi & Sami Belkacem

Exploratory Data Analysis (EDA)

- ❑ EDA is a set of **statistical** and **visualization** techniques.

- ❑ Used for seeing what the data can tell before the **preprocessing** and **modeling**:
 - ❑ Understand the data and summarize its keys properties.
 - ❑ Discover noisy data and outliers.
 - ❑ Comprehend the distribution of the data.
 - ❑ Decide which data cleaning techniques to be applied.

- ❑ EDA is cross-classified in two ways:
 - ❑ The method is either **non-graphical** or **graphical**.
 - ❑ The method is either **univariate** or **multivariate** (usually just bivariate).

Exploratory Data Analysis (EDA)

1. Statics of Data
2. Data visualization
3. How to choose an EDA technique


1. Statics of Data

- ❑ Population vs Sample
- ❑ Measuring the Central Tendency (Mean, Median, and Mode)
- ❑ Symmetric vs. Skewed Data
- ❑ Measuring data distribution (Variance and Standard deviation)
- ❑ Analysis of two variables (Covariance and Correlation)

Population vs Sample

- ❑ **Population:** The entire group that we want to draw conclusions about.
- ❑ **Sample:** Subset of the population, used when the population size is too large.

It should be an unbiased subset that best represents the entire population.



Measuring the Central Tendency: (1) Mean

- Mean (algebraic measure) (sample vs. population):

n is sample size and N is population size.

$$\bar{x} = \frac{1}{n} \sum_{i=1}^n x_i \quad \mu = \frac{\sum x}{N}$$

- **Weighted arithmetic mean:**

$$\bar{x} = \frac{\sum_{i=1}^n w_i x_i}{\sum_{i=1}^n w_i}$$

- **Trimmed mean:** Chopping extreme values (e.g., olympics gymnastics score computation)

Measuring the Central Tendency: (2) Median

- ❑ **Median** is the middle value in a data set when values are ordered.

- ❑ **How to calculate**
 - ❑ Sort data in ascending order.
 - ❑ Repeat values according to their frequency.
 - ❑ Odd number of data points => the median is the middle value.
 - ❑ Even number of data points => the median is the average of the two middle values.

- ❑ **Why use the Median**
 - ❑ Resistant to extreme outliers compared to the mean.
 - ❑ Useful for skewed distributions.

Median calculation for grouped data

$$\text{Estimated Median} = L + \frac{(n/2) - B}{G} \times w$$

- ❑ **L**: lower class boundary of the median bin.
- ❑ **n**: total number of values.
- ❑ **B**: cumulative frequency of the bins before the median bin.
- ❑ **G**: frequency of the median bin.
- ❑ **W**: median bin width.

	Age	Frequency
0	5-10	4
1	10-20	5
2	20-50	3
3	50-55	1

Median calculation for grouped data

$$\text{Estimated Median} = L + \frac{(n/2) - B}{G} \times w$$

	Age	Frequency
0	5-10	4
1	10-20	5
2	20-50	3
3	50-55	1

❑ The median bin is the value 10-20

- ❑ $L = 10$
- ❑ $n = 13$
- ❑ $B = 4$
- ❑ $G = 5$
- ❑ $w = 20 - 10 = 10$


$$\text{Median value} = 10 + ((6.5 - 4) / 5) * 10 = 15$$

Measuring the Central Tendency: (3) Mode

□ **Mode:** Value that occurs most frequently in the data

□ **Unimodal**


□ Empirical formula: $mean - mode = 3 \times (mean - median)$




□ **Multimodal**

□ Bimodal (a)

□ Trimodal (b)




(a)




(b)

Symmetric vs. Skewed Data


- Median, mean, and mode of symmetric, positively, and negatively skewed data



Positively Skewed




Symmetric
(Normal distribution)



Negatively Skewed


Properties of Normal Distribution Curve



Represent central tendency 12

Properties of Data after z-score normalization

Z-score normalization (μ : mean, σ : standard deviation):
$$z = \frac{X - \mu}{\sigma}$$



Measuring Data Distribution: Variance and Standard Deviation

□ *s*: sample, *σ*: population, *n*: size of the sample, *N*: size of the population

□ **Variance** : Measure dispersion around the mean

$$s^2 = \frac{1}{n-1} \sum_{i=1}^n (x_i - \bar{x})^2 = \frac{1}{n-1} \left[\sum_{i=1}^n x_i^2 - \frac{1}{n} \left(\sum_{i=1}^n x_i \right)^2 \right]$$

$$\sigma^2 = \frac{1}{N} \sum_{i=1}^n (x_i - \mu)^2 = \frac{1}{N} \sum_{i=1}^n x_i^2 - \mu^2$$

□ For a **sample**: Divide by *n-1* to account for bias in the estimation.

For a **population**: Divide by *N* because no correction is needed.

□ **Standard deviation** *s* (or *σ*) is the square root of variance *s*² (or *σ*²)

Measures dispersion around the mean, but in the same units as the value

Analysis of two variables: Covariance

$$cov_{x,y} = \frac{\sum (x_i - \bar{x})(y_i - \bar{y})}{N - 1}$$

- ❑ The covariance for two numerical variables computes the relationship between them.
- ❑ Covariance shows how two variables change together:
 - ❑ **Positive Covariance:** both variables move together.
 - ❑ **Negative Covariance:** variables move in opposite directions.
 - ❑ **Zero Covariance:** no clear pattern in variable movements.
- ❑ Covariance is sensitive to scale, correlation addresses this by standardizing the measurement.

Covariance Matrix

- The variance and covariance information for the two variables can be summarized as 2X2 covariance matrix, which contains the following information:
 - The variance of each variable on the diagonal.
 - The covariance between the two variables in the off-diagonal entries.

$$\Sigma_{i,j} = \begin{bmatrix} \sigma_{ii}^2 & \sigma_{ij}^2 \\ \sigma_{ji}^2 & \sigma_{jj}^2 \end{bmatrix}$$

- Generalizing it to d dimensions: $\Sigma = E[(\mathbf{X} - \mu)(\mathbf{X} - \mu)^T] = \begin{pmatrix} \sigma_1^2 & \sigma_{12} & \cdots & \sigma_{1d} \\ \sigma_{21} & \sigma_2^2 & \cdots & \sigma_{2d} \\ \vdots & \vdots & \ddots & \vdots \\ \sigma_{d1} & \sigma_{d2} & \cdots & \sigma_d^2 \end{pmatrix}$

Analysis of two variables: Correlation

- **Correlation** between two numerical variables $\mathbf{X}_1$ and $\mathbf{X}_2$ is the standard covariance, obtained by normalizing the covariance with the standard deviation of each variable.

$$\rho_{12} = \frac{\sigma_{12}}{\sigma_1 \sigma_2} = \frac{\sigma_{12}}{\sqrt{\sigma_1^2 \sigma_2^2}}$$


- **Sample correlation** for two attributes X_1 and X_2 : $\hat{\rho}_{12} = \frac{\hat{\sigma}_{12}}{\hat{\sigma}_1 \hat{\sigma}_2} = \frac{\sum_{i=1}^n (x_{i1} - \hat{\mu}_1)(x_{i2} - \hat{\mu}_2)}{\sqrt{\sum_{i=1}^n (x_{i1} - \hat{\mu}_1)^2 \sum_{i=1}^n (x_{i2} - \hat{\mu}_2)^2}}$

n : number of tuples, μ_1 and μ_2 : means of $\mathbf{X}_1$ and $\mathbf{X}_2$, σ_1 and σ_2 : standard deviation of $\mathbf{X}_1$ and $\mathbf{X}_2$


- **If $\rho_{12} > 0$:** A and B are positively correlated (X_1 's values increase as X_2 's)
- **If $\rho_{12} < 0$:** A and B are negatively correlated (X_1 's values increases, X_2 's decreases)
- **If $\rho_{12} = 0$:** independent (under the same assumption as discussed in co-variance)

Correlation Matrix (Correlation Heatmap)

- Matrix that shows the correlations between each pair of variables in a dataset



Visualizing Changes of Correlation Coefficient




- Correlation coefficient value range: $[-1, 1]$
- A set of scatter plots shows sets of points and their correlation coefficients changing from -1 to 1

2. Data Visualization


- ☐ **Boxplot**
- ☐ **Histogram and Bar chart**
- ☐ **Quantile plot**
- ☐ **Quantile-quantile (Q-Q) plot**
- ☐ **Scatter plot**
- ☐ **Line chart**
- ☐ **Parallel Coordinates plot**

Boxplot: Data Dispersion

- ❑ **Quartiles:** Q_1 (25th percentile), Q_3 (75th percentile)
- ❑ **Interquartile range:** $IQR = Q_3 - Q_1$
- ❑ **Five number summary:** min, Q_1 , median, Q_3 , max
- ❑ **Box plot:** Data is represented with a box
 - ❑ **Q_1 , Q_3 , IQR:** The ends of the box are at the first and third quartiles, i.e., the height of the box is IQR
 - ❑ **Median (Q_2)** is marked by a line within the box
 - ❑ **Whiskers:** two lines outside the box extended to Minimum and Maximum
 - ❑ **Outliers:** points beyond a specified threshold (e.g. value higher/lower than $1.5 \times IQR$)





Boxplot: detect outliers





Histogram and Bar chart

- **Histogram:** Tabulated frequencies represented by bars.
- **Bar chart:** Categorical data with bars proportional to the values they represent.





Histogram and Bar chart

- ❑ Histograms show **distributions** of variables, while bar charts compare **variables**
- ❑ Histograms plot binned **quantitative/categorical** data, while bar charts only plot **categorical** data
- ❑ Bars can be **reordered** in bar charts, but not in histograms
- ❑ In histograms, it is the area of the bar that denotes the value, not the height as in bar charts.




Histogram Analysis

- Divide data into buckets and store average (sum) for each bucket.
- Partitioning rules:
 - **Equal-width:** equal bucket range
 - **Equal-frequency** (or equal-depth)




Histograms often tell more than Boxplots



- The two histograms shown in the left may have the same box plot representation

- The same values for: min, Q1, median, Q3, max




- But they have rather different data distributions

Quantile Plot

- ❑ **Purpose:** Visualizes all quantile information for a specific attribute.

- ❑ **Benefits:**

- ❑ Provides a comprehensive view of the attribute's distribution
- ❑ Helps identify general trends and outliers



- ❑ **Construction:**

- ❑ For a data $\{X_1, X_2, \dots, X_N\}$ sorted in increasing order.
- ❑ f_i indicates that approximately f_i of the data point have values $\leq x_i$

Quantile-Quantile (Q-Q) Plot

Definition:

Graphical tool used to compare the quantiles of two distributions.



- **X-axis:** Quantiles of the theoretical distribution
- **Y-axis:** Quantiles of the data distribution
- **Reference line:** A diagonal line (usually $y=x$)

Purpose:

Assess the distributional similarity between an attribute and **another attribute** or a **theoretical distribution (e.g. normal distribution)**.



Interpretation:

- **Close fit to straight Line:** The two distributions are similar
- **Deviations from the Line:** The two distributions are different




Scatter plot

- ❑ Provides a first look at bivariate data to see clusters of points, outliers, etc.
- ❑ Each pair of values is treated as a pair of coordinates and plotted as points in the plane.




Line chart

- ❑ A line chart displays information as a series of data points called 'markers'.
- ❑ The markers are connected to each other by straight line segments.



Parallel Coordinates Plots

- A parallel coordinate plot maps each object in the dataset as a line.
- Each attribute of an object is represented by a point on the line.
- The following example shows the parallel coordinates plots of the iris dataset.



3. How to choose an EDA technique

Criteria to choose a statistical/visualization technique:

- ❑ The type of analysis:
 - ❑ comparison, relationship, composition, distribution, dispersion, etc.
- ❑ The number of attributes (one, two, multiple attributes, etc.)
- ❑ The types of attributes (categorical, numerical, etc.)

The summary statistics: [Summary statistics.pdf](#)

The chart chooser: [The chart chooser.pdf](#)

Understanding Question

Given a dataset of objects, how to visually explore the relationship between two variables in the following cases?

- Continuous vs Continuous
- Categorical vs Continuous
- Continuous vs Categorical
- Categorical vs Categorical