

DATABASES

Chapter 7 : Database Performance Tuning and Query Optimization

Pr. Kamel BOUKHALFA

Slides From the Textbook :
Carlos Coronel and Steven Morris, Database Systems: Design, Implementation, and Management
Tenth Edition

Objectives

- In this chapter, you will learn:
 - Basic database performance-tuning concepts
 - How a DBMS processes SQL queries
 - About the importance of indexes in query processing

Objectives (continued)



- In this chapter, you will learn: (continued)
 - About the types of decisions the query optimizer has to make
 - Some common practices used to write efficient SQL code
 - How to formulate queries and tune the DBMS for optimal performance

3

Database Performance-Tuning Concepts



- Goal of database performance is to execute queries as fast as possible
- **Database performance tuning**
 - Set of activities and procedures designed to reduce response time of database system
- All factors must operate at optimum level with minimal bottlenecks
- **Good database performance starts with good database design**

4

TABLE
11.1

General Guidelines for Better System Performance

	SYSTEM RESOURCES	CLIENT	SERVER
Hardware	CPU	The fastest possible Dual-core CPU or higher	The fastest possible Multiple processors (Quad-core technology)
	RAM	The maximum possible	The maximum possible
	Hard Disk	Fast SATA/EIDE hard disk with sufficient free hard disk space	Multiple high-speed, high-capacity hard disks (SCSI / SATA / Firewire / Fibre Channel) in RAID configuration
	Network	High-speed connection	High-speed connection
Software	Operating System	Fine-tuned for best client application performance	Fine-tuned for best server application performance
	Network	Fine-tuned for best throughput	Fine-tuned for best throughput
	Application	Optimize SQL in client application	Optimize DBMS server for best performance

5

Performance Tuning: Client and Server

- Client side
 - Generate SQL query that returns correct answer in least amount of time
 - Using minimum amount of resources at server
 - **SQL performance tuning**
- Server side
 - DBMS environment configured to respond to clients' requests as fast as possible
 - Optimum use of existing resources
 - **DBMS performance tuning**

6


DBMS Architecture



- All data in database are stored in data files
- **Data files**
 - Automatically expand in predefined increments known as **extends**
 - Grouped in **file groups** or **table spaces**
- Table space or file group:
 - Logical grouping of several data files that store data with similar characteristics

7

11.1



8

DBMS Architecture (continued)



- **Data cache** or **buffer cache**: shared, reserved memory area
 - Stores most recently accessed data blocks in RAM
- **SQL cache** or **procedure cache**: stores most recently executed SQL statements
 - Also PL/SQL procedures, including triggers and functions
- DBMS retrieves data from permanent storage and places it in RAM

9

DBMS Architecture (continued)



- **Input/output request**: low-level data access operation to/from computer devices
- Data cache is faster than data in data files
 - DBMS does not wait for hard disk to retrieve data
- Majority of performance-tuning activities focus on minimizing I/O operations
- Typical DBMS processes:
 - Listener, User, Scheduler, Lock manager, Optimizer

10

Database Statistics



- Measurements about database objects and available resources
 - Tables
 - Indexes
 - Number of processors used
 - Processor speed
 - Temporary space available

11

Database Statistics (continued)



- Make critical decisions about improving query processing efficiency
- Can be gathered manually by DBA or automatically by DBMS

12

TABLE
11.2

Sample Database Statistics Measurements


DATABASE OBJECT	SAMPLE MEASUREMENTS
Tables	Number of rows, number of disk blocks used, row length, number of columns in each row, number of distinct values in each column, maximum value in each column, minimum value in each column, and columns that have indexes
Indexes	Number and name of columns in the index key, number of key values in the index, number of distinct key values in the index key, histogram of key values in an index, and number of disk pages used by the index
Environment Resources	Logical and physical disk block size, location and size of data files, and number of extents per data file

13

Query Processing

- DBMS processes queries in three phases
 - Parsing
 - DBMS parses the query and chooses the most efficient access/execution plan
 - Execution
 - DBMS executes the query using chosen execution plan
 - Fetching
 - DBMS fetches the data and sends the result back to the client

14



15

SQL Parsing Phase

- Break down query into smaller units
- Transform original SQL query into slightly different version of original SQL code
 - Fully equivalent
 - Optimized query results are always the same as original query
 - More efficient
 - Optimized query will almost always execute faster than original query

16

SQL Parsing Phase (continued)



- Query optimizer analyzes SQL query and finds most efficient way to access data
 - Validated for syntax compliance
 - Validated against data dictionary
 - Tables, column names are correct
 - User has proper access rights
 - Analyzed and decomposed into components
 - Optimized
 - Prepared for execution

17

SQL Parsing Phase (continued)



- Access plans are DBMS-specific
 - Translate client's SQL query into series of complex I/O operations
 - Required to read the data from the physical data files and generate result set
- DBMS checks if access plan already exists for query in SQL cache
- DBMS reuses the access plan to save time
- If not, optimizer evaluates various plans
 - Chosen plan placed in SQL cache

18

TABLE
11.3

Sample DBMS Access Plan I/O Operations

OPERATION	DESCRIPTION
Table Scan (Full)	Reads the entire table sequentially, from the first row to the last row, one row at a time (slowest)
Table Access (Row ID)	Reads a table row directly, using the row ID value (fastest)
Index Scan (Range)	Reads the index first to obtain the row IDs and then accesses the table rows directly (faster than a full table scan)
Index Access (Unique)	Used when a table has a unique index in a column
Nested Loop	Reads and compares a set of values to another set of values, using a nested loop style (slow)
Merge	Merges two data sets (slow)
Sort	Sorts a data set (slow)

19

SQL Execution Phase

SQL Fetching Phase

- All I/O operations indicated in access plan are executed
 - Locks acquired
 - Data retrieved and placed in data cache
 - Transaction management commands processed
- Rows of resulting query result set are returned to client
- DBMS may use temporary table space to store temporary data

20

Query Processing Bottlenecks



- Delay introduced in the processing of an I/O operation that slows the system
 - CPU
 - RAM
 - Hard disk
 - Network
 - Application code

21

Indexes and Query Optimization



- Indexes
 - Crucial in speeding up data access
 - Facilitate searching, sorting, and using aggregate functions as well as join operations
 - Ordered set of values that contains index key and pointers
- More efficient to use index to access table than to scan all rows in table sequentially

22

Indexes and Query Optimization (continued) ensia

- **Data sparsity:** number of different values a column could possibly have
- Indexes implemented using:
 - Hash indexes
 - B-tree indexes
 - Bitmap indexes
- DBMSs determine best type of index to use

23

FIGURE 11.3 Index representation for the CUSTOMER table

STATE_NDX INDEX

Key	Row
AZ	2
...	...
FL	1
...	...
FL	7
...	...
FL	9
...	...
FL	13245
...	...
FL	14786
...	...

CUSTOMER TABLE
(14,786 rows)

Row ID	CUS_CODE	CUS_LNAME	CUS_FNAME	CUS_INITIAL	CUS_AREACODE	CUS_PHONE	CUS_STATE	CUS_BALANCE
1	10010	Farniss	Alfred	A	815	944-2513	FL	\$0.00
2	10011	Dunne	Leona	K	713	894-1238	AZ	\$0.00
3	10012	Smith	Kathy	W	815	894-2285	TX	\$345.86
4	10013	Olowski	Paul	F	815	894-2180	AZ	\$536.75
5	10014	Orlondo	Myron		815	222-1872	NY	\$0.00
6	10015	O'Brian	Amy	B	713	442-3381	NY	\$0.00
7	10016	Brown	James	G	815	297-1228	FL	\$221.19
8	10017	Williams	George		815	290-2556	FL	\$768.93
9	10018	Farniss	Anne	G	713	382-7185	TX	\$216.55
10	10019	Smith	Olette	K	815	297-3809	AZ	\$0.00
...	...	...	...	...	...	...	...	...
13245	23120	Veron	George	D	415	231-9872	FL	\$875.00
...	...	...	...	...	...	...	...	...
14786	24560	Suarez	Victor		435	342-9876	FL	\$342.00

24

Optimizer Choices



- **Rule-based optimizer**
 - Preset rules and points
 - Rules assign a fixed cost to each operation
- **Cost-based optimizer**
 - Algorithms based on statistics about objects being accessed
 - Adds up processing cost, I/O costs, resource costs to derive total cost

25



TABLE 11.4 Comparing Access Plans and I/O Costs

PLAN	STEP	OPERATION	I/O OPERATIONS	I/O COST	RESULTING SET ROWS	TOTAL I/O COST
A	A1	Cartesian product (PRODUCT, VENDOR)	7,000 + 300	7,300	2,100,000	7,300
	A2	Select rows in A1 with matching vendor codes	2,100,000	2,100,000	7,000	2,107,300
	A3	Select rows in A2 with V_STATE = 'FL'	7,000	7,000	1,000	2,114,300
B	B1	Select rows in VENDOR with V_STATE = 'FL'	300	300	10	300
	B2	Cartesian Product (PRODUCT, B1)	7,000 + 10	7,010	70,000	7,310
	B3	Select rows in B2 with matching vendor codes	70,000	70,000	1,000	77,310

26

Using Hints to Affect Optimizer Choices



- Optimizer might not choose best plan
- Makes decisions based on existing statistics
 - Statistics may be old
 - Might choose less efficient decisions
- **Optimizer hints:** special instructions for the optimizer embedded in the SQL command text

27

TABLE
11.5

Optimizer Hints

HINT	USAGE
ALL_ROWS	Instructs the optimizer to minimize the overall execution time, that is, to minimize the time it takes to return all rows in the query result set. This hint is generally used for batch mode processes. For example: <pre>SELECT /*+ ALL_ROWS */ * FROM PRODUCT WHERE P_QOH < 10;</pre>
FIRST_ROWS	Instructs the optimizer to minimize the time it takes to process the first set of rows, that is, to minimize the time it takes to return only the first set of rows in the query result set. This hint is generally used for interactive mode processes. For example: <pre>SELECT /*+ FIRST_ROWS */ * FROM PRODUCT WHERE P_QOH < 10;</pre>
INDEX(name)	Forces the optimizer to use the P_QOH_NDX index to process this query. For example: <pre>SELECT /*+ INDEX(P_QOH_NDX) */ * FROM PRODUCT WHERE P_QOH < 10;</pre>

28

SQL Performance Tuning



- Evaluated from client perspective
 - Most current relational DBMSs perform automatic query optimization at the server end
 - Most SQL performance optimization techniques are DBMS-specific
 - Rarely portable
- Majority of performance problems related to poorly written SQL code
- Carefully written query usually outperforms a poorly written query

29

Index Selectivity



- Indexes are used when:
 - Indexed column appears by itself in search criteria of WHERE or HAVING clause
 - Indexed column appears by itself in GROUP BY or ORDER BY clause
 - MAX or MIN function is applied to indexed column
 - Data sparsity on indexed column is high
- Measure of how likely an index will be used

30

Index Selectivity (continued)



- General guidelines for indexes:
 - Create indexes for each attribute in WHERE, HAVING, ORDER BY, or GROUP BY clause
 - Do not use in small tables or tables with low sparsity
 - Declare primary and foreign keys so optimizer can use indexes in join operations
 - Declare indexes in join columns other than PK/FK

31

Conditional Expressions



- Normally expressed within WHERE or HAVING clauses of SQL statement
- Restricts output of query to only rows matching conditional criteria

32

Conditional Expressions (continued)



- Common practices for efficient SQL:
 - Use simple columns or literals in conditionals
 - Numeric field comparisons are faster
 - Equality comparisons faster than inequality
 - Transform conditional expressions to use literals
 - Write equality conditions first
 - AND: Use condition most likely to be false first
 - OR: Use condition most likely to be true first
 - Avoid NOT

33

Query Formulation



- Identify what columns and computations are required
- Identify source tables
- Determine how to join tables
- Determine what selection criteria is needed
- Determine in what order to display output

34

DBMS Performance Tuning



- Includes managing DBMS processes in primary memory and structures in physical storage
- DBMS performance tuning at server end focuses on setting parameters used for:
 - Data cache
 - SQL cache
 - Sort cache
 - Optimizer mode

35

DBMS Performance Tuning (continued)



- Some general recommendations for creation of databases:
 - Use **RAID** (Redundant Array of Independent Disks) to provide balance between performance and fault tolerance
 - Minimize disk contention
 - Put high-usage tables in their own table spaces
 - Assign separate data files in separate storage volumes for indexes, system, high-usage tables

36

DBMS Performance Tuning (continued)

- Some general recommendations for creation of databases: (continued)
 - Take advantage of table storage organizations in database
 - Partition tables based on usage
 - Use denormalized tables where appropriate
 - Store computed and aggregate attributes in tables

37

Query Optimization Example

- Example illustrates how query optimizer works
- Based on QOVENDOR and QOPRODUCT tables
- Uses Oracle SQL*Plus

38

11.4


B-Tree index is used in columns with high data sparsity – that is, columns with many different values relative to the total number of rows.

CUSTOMER TABLE

CUS_ID	CUS_LNAME	CUS_FNAME	CUS_PHONE	REGION_CODE
12	Adams	Charles	4353	NW
23	Blair	Robert	5426	SE
37	Connell	Charles	2398	SW
43	DePuccio	Robert	6643	NE
55	Greer	Tim	2764	SE
59	Fule	Ruben	2453	SW
62	Lee	John	7696	NE
65	Mason	Jerry	7609	NW
69	Morris	Steve	4468	NW
72	Rob	Pete	8123	NE
75	Server	Lee	8193	SE
80	Stickland	Thomas	3129	SW
82	Tennison	Douglas	3499	NE

Bitmap index is used in columns with low data sparsity – that is, columns with few different values relative to the total number of rows.

B-tree Index
On CUS_LNAME



Leaf objects contain index key and pointers to rows in table. Access to any row using the index will take the same number of I/O accesses. In this example it would take 4 I/O accesses to access any given table row using the index. One for each index tree level (root, branch, leaf object) plus access to data row using the pointer.

Bitmap Index
On REGION_CODE

Region →

	NE	NW	SE	SW	...
Bit 1	0	1	0	0	...
Bit 2	0	0	1	0	...
Bit 3	0	0	0	1	...
Bit 4	1	0	0	0	...
Bit 5	0	0	1	0	...
Bit 6	1	0	0	0	...
Bit 7	0	1	0	0	...
Bit 8	0	1	0	0	...
Bit 9	1	0	0	0	...
Bit 10	0	0	1	0	...
Bit 11	0	0	0	1	...
Bit 12	1	0	0	0	...

← One byte

In the bitmap index, each bit represents one region code. In the first row, bit number two is turned on, thus indicating that the first row region code value is NW.

REGION_CODE = 'NW'

Each byte in the bitmap index represents one row of the table data. Bitmap indexes are very efficient with searches. For example, to find all customers in the NW region, the DBMS will return all rows with bit number two turned on.

39

11.5

```

Oracle SQL*Plus
File Edit Search Options Help
SQL> ANALYZE TABLE QOVENDOR COMPUTE STATISTICS;
Table analyzed.
SQL> EXPLAIN PLAN FOR SELECT * FROM QOVENDOR WHERE U_NAME LIKE 'B%' ORDER BY U_AREACODE;
Explained.
SQL> SELECT * FROM TABLE(DBMS_XPLAN.DISPLAY);
PLAN_TABLE_OUTPUT
Plan hash value: 1837703589

   Id | Operation              | Name       | Rows  | Bytes | Cost (%CPU)| Time     |
----|-----|-----|-----|-----|-----|-----|-----|
   0 | SELECT STATEMENT        |            |       1 |    38 |  4 (25)| 00:00:01 |
   1 |   SORT ORDER BY         |            |       1 |    38 |  4 (25)| 00:00:01 |
   2 |    TABLE ACCESS FULL   | QOVENDOR   |       1 |    38 |  3 (0)| 00:00:01 |

Predicate Information (identified by operation id):
PLAN_TABLE_OUTPUT

      2 - filter("U_NAME" LIKE 'B%')

14 rows selected.
SQL>
  
```

40

11.6

```

Oracle SQL*Plus
File Edit Search Options Help
SQL> CREATE INDEX QOV_NDX1 ON QOVENDOR(U_AREACODE);
Index created.
SQL> ANALYZE TABLE QOVENDOR COMPUTE STATISTICS;
Table analyzed.
SQL> EXPLAIN PLAN FOR SELECT * FROM QOVENDOR WHERE U_NAME LIKE 'B%' ORDER BY U_AREACODE;
Explained.
SQL> SELECT * FROM TABLE(DBMS_XPLAN.DISPLAY);
PLAN_TABLE_OUTPUT
Plan hash value: 2305289760

   Id  Operation                Name          Rows  Bytes  Cost (%CPU)  Time
--  -  -
   0   SELECT STATEMENT
   *  1   TABLE ACCESS BY INDEX ROWID QOVENDOR         1     38       2   (0)   00:00:01
   2   INDEX FULL SCAN          QOV_NDX1      15              1   (0)   00:00:01

Predicate Information (identified by operation id):
   1 - filter("U_NAME" LIKE 'B%')

14 rows selected.

SQL>

```

41

11.7

```

Oracle SQL*Plus
File Edit Search Options Help
SQL> CREATE INDEX QOV_NDX2 ON QOVENDOR(U_NAME);
Index created.
SQL> ANALYZE TABLE QOVENDOR COMPUTE STATISTICS;
Table analyzed.
SQL> EXPLAIN PLAN FOR SELECT * FROM QOVENDOR WHERE U_NAME LIKE 'B%' ORDER BY U_AREACODE;
Explained.
SQL> SELECT * FROM TABLE(DBMS_XPLAN.DISPLAY);
PLAN_TABLE_OUTPUT
Plan hash value: 2305289760

   Id  Operation                Name          Rows  Bytes  Cost (%CPU)  Time
--  -  -
   0   SELECT STATEMENT
   *  1   TABLE ACCESS BY INDEX ROWID QOVENDOR         1     38       2   (0)   00:00:01
   2   INDEX FULL SCAN          QOV_NDX1      15              1   (0)   00:00:01

Predicate Information (identified by operation id):
   1 - filter("U_NAME" LIKE 'B%')

14 rows selected.

SQL>

```

42

11.8

```

Oracle SQL*Plus
File Edit Search Options Help
SQL> EXPLAIN PLAN FOR SELECT U_NAME, P_CODE FROM QOVENDOR U, QOPRODUCT P
2 WHERE U.U_CODE = P.U_CODE AND U_NAME = 'ORDUA, Inc.';

Explained.

SQL> SELECT * FROM TABLE(DBMS_XPLAN.DISPLAY);

PLAN_TABLE_OUTPUT
-----
Plan hash value: 3956542569

   Id | Operation              | Name              | Rows  | Bytes | Cost (%CPU)| Time     |
---- |-----|-----|-----|-----|-----|-----|-----|
| 0 | SELECT STATEMENT        |                   |      2 |    74 | 6 (17)| 00:00:01 |
|* 1 | HASH JOIN               |                   |      2 |    74 | 6 (17)| 00:00:01 |
| 2 | TABLE ACCESS BY INDEX ROWID | QOVENDOR          |      1 |    17 | 2 (0)| 00:00:01 |
|* 3 | INDEX RANGE SCAN        | QOV_MIDX2         |      1 |      | 1 (0)| 00:00:01 |
| 4 | TABLE ACCESS FULL      | QOPRODUCT         |     16 |   320 | 3 (0)| 00:00:01 |

Predicate Information (identified by operation id):

   1 - access("U"."U_CODE"="P"."U_CODE")
   3 - access("U_NAME"='ORDUA, Inc.')

Note
-----
   - dynamic sampling used for this statement

21 rows selected.

SQL>

```

43

11.9

```

Oracle SQL*Plus
File Edit Search Options Help
SQL> EXPLAIN PLAN FOR SELECT U_NAME, P_CODE FROM QOVENDOR U, QOPRODUCT P
2 WHERE U.U_CODE = P.U_CODE AND UPPER(U_NAME) = 'ORDUA, INC.';

Explained.

SQL> SELECT * FROM TABLE(DBMS_XPLAN.DISPLAY);

PLAN_TABLE_OUTPUT
-----
Plan hash value: 4061476548

   Id | Operation              | Name              | Rows  | Bytes | Cost (%CPU)| Time     |
---- |-----|-----|-----|-----|-----|-----|-----|
| 0 | SELECT STATEMENT        |                   |      1 |    37 | 7 (15)| 00:00:01 |
|* 1 | HASH JOIN               |                   |      1 |    37 | 7 (15)| 00:00:01 |
| 2 | VIEW                    | index$_join$_001  |      1 |    17 | 3 (0)| 00:00:01 |
|* 3 | HASH JOIN               |                   |      1 |    17 | 1 (0)| 00:00:01 |
| 4 | INDEX FAST FULL SCAN    | QOV_MIDX2         |      1 |    17 | 1 (0)| 00:00:01 |
| 5 | INDEX FAST FULL SCAN    | SYS_C006002       |      1 |    17 | 1 (0)| 00:00:01 |
| 6 | TABLE ACCESS FULL      | QOPRODUCT         |     16 |   320 | 3 (0)| 00:00:01 |

Predicate Information (identified by operation id):

   1 - access("U"."U_CODE"="P"."U_CODE")
   2 - filter(UPPER("U_NAME")='ORDUA, INC.')
   3 - access(ROWID=ROWID)
   4 - filter(UPPER("U_NAME")='ORDUA, INC.')

Note
-----
   - dynamic sampling used for this statement

25 rows selected.

SQL>

```

44

11.10

```

Oracle SQL*Plus
File Edit Search Options Help
SQL> ANALYZE TABLE QOPRODUCT COMPUTE STATISTICS;

Table analyzed.

SQL> EXPLAIN PLAN FOR SELECT MAX(P_PRICE) FROM QOPRODUCT;

Explained.

SQL> SELECT * FROM TABLE(DBMS_XPLAN.DISPLAY);

PLAN_TABLE_OUTPUT
-----
Plan hash value: 2544502246

   Id | Operation          | Name                | Rows | Bytes | Cost (%CPU)| Time     |
----|-----|-----|-----|-----|-----|-----|-----|
   0 | SELECT STATEMENT    |                     |      |       |            |          |
   1 |  SORT AGGREGATE     |                     |      |       |            |          |
   2 |   TABLE ACCESS FULL| QOPRODUCT           |    16 |    64 |      3 (0)| 00:00:01 |

9 rows selected.

SQL>

```

45

11.11

```

Oracle SQL*Plus
File Edit Search Options Help
SQL> CREATE INDEX QOP_IDX2 ON QOPRODUCT(P_PRICE);

Index created.

SQL> ANALYZE TABLE QOPRODUCT COMPUTE STATISTICS;

Table analyzed.

SQL> EXPLAIN PLAN FOR SELECT MAX(P_PRICE) FROM QOPRODUCT;

Explained.

SQL> SELECT * FROM TABLE(DBMS_XPLAN.DISPLAY);

PLAN_TABLE_OUTPUT
-----
Plan hash value: 3423609809


   Id | Operation                      | Name                | Rows | Bytes | Cost (%CPU)| Time     |
----|-----|-----|-----|-----|-----|-----|-----|
   0 | SELECT STATEMENT                |                     |      |       |            |          |
   1 |  SORT AGGREGATE                 |                     |      |       |            |          |
   2 |   INDEX FULL SCAN (MIN/MAX)     | QOP_IDX2            |    16 |    64 |      1 (0)| 00:00:01 |

9 rows selected.

SQL>

```

46



47

Summary

- Database performance tuning
 - Refers to activities to ensure query is processed in minimum amount of time
- SQL performance tuning
 - Refers to activities on client side to generate SQL code
 - Returns correct answer in least amount of time
 - Uses minimum amount of resources at server end
- DBMS architecture represented by processes and structures used to manage a database

48

Summary (continued)



- Database statistics refers to measurements gathered by the DBMS
 - Describe snapshot of database objects' characteristics
- DBMS processes queries in three phases: parsing, execution, and fetching
- Indexes are crucial in process that speeds up data access

49

Summary (continued)



- During query optimization, DBMS chooses:
 - Indexes to use, how to perform join operations, table to use first, etc.
- Hints change optimizer mode for current SQL statement
- SQL performance tuning deals with writing queries that make good use of statistics
- Query formulation deals with translating business questions into specific SQL code

50