

Privacy-Preserving Face Recognition Using Federated Learning: A Comprehensive Research Report

Introduction to the Decentralized Biometric Paradigm

Facial recognition technology has experienced a rapid and transformative evolution over the past decade, transitioning from foundational computer vision algorithms reliant on handcrafted features to highly sophisticated deep learning architectures. Modern artificial intelligence frameworks now routinely surpass human-level performance on diverse and complex datasets.¹ These profound technological advancements have facilitated the ubiquitous integration of facial recognition across a multitude of critical sectors, ranging from expansive security monitoring systems and border control infrastructure to everyday mobile device authentication, cell phone unlocking mechanisms, and enterprise-level online identity verification.² However, the proliferation of centralized biometric databases—which function as the traditional backbone for training these deep learning models—has precipitated severe and unprecedented privacy and security vulnerabilities. Facial data is inherently immutable; unlike a compromised password or cryptographic key, biometric identifiers cannot be revoked or reissued. They contain highly sensitive personally identifiable information that is permanently linked to an individual's physical identity.¹

The catastrophic risks associated with centralized data repositories have been repeatedly underscored by high-profile data breaches. For instance, the Biostar2 incident resulted in the leak of over one million biometric records, encompassing both fingerprint and facial recognition information, alongside unencrypted administrative credentials managed by a security firm serving the banking sector and government contractors.³ Similarly, unauthorized access to the genetic matching platform of 23andMe compromised the highly sensitive DNA, profile, and ethnicity information of nearly seven million individuals.³ In other instances, stolen identity credentials have been weaponized in automated injection attacks, such as credential stuffing, targeting government infrastructure like the Government of Canada.³ Simultaneously, the regulatory landscape governing data privacy has become increasingly stringent. Frameworks such as the European Union's General Data Protection Regulation (GDPR), the California Consumer Privacy Act (CCPA), and the Health Insurance Portability and Accountability Act (HIPAA) impose severe legal, financial, and operational constraints on the collection, centralized storage, and cross-border transfer of biometric and health-related data.⁴ These developments have created a critical paradox for the artificial intelligence community: the demand for high-accuracy, data-hungry facial recognition models continues to accelerate, while the legal and ethical mandates to protect the underlying training data have

never been stricter.

To reconcile this tension, researchers have increasingly pivoted toward Federated Learning (FL). Federated learning represents an innovative, paradigm-shifting decentralized machine learning architecture that enables multiple distributed edge devices, mobile endpoints, or institutional entities to collaboratively train a shared global model without ever transferring raw, localized data to a central server.¹ By transmitting only computed model parameters or mathematical gradients, FL intrinsically mitigates the risk of massive data exposure, ensuring enhanced privacy and security for each participating node.¹ However, the integration of federated learning into facial recognition is not a frictionless, plug-and-play solution. The decentralized architecture introduces a host of novel, complex systemic vulnerabilities. These include sophisticated gradient-based privacy attacks capable of reconstructing training data, Byzantine poisoning behaviors that threaten model integrity, and severe optimization challenges stemming from non-independent and identically distributed (non-IID) data distributions across highly heterogeneous edge devices.²

The following exhaustive analysis provides a comprehensive examination of the intersection between facial recognition and federated learning spanning the 2024 to 2026 research continuum. It dissects the underlying mathematical architectures, security vulnerabilities, advanced cryptographic enhancements, state-of-the-art algorithmic innovations (such as FedGC and VOIDFace), and the profound hardware constraints that define the current landscape of decentralized biometric authentication.

Architectural and Mathematical Foundations of Federated Optimization

Traditional machine learning frameworks rely on a centralized paradigm where immense volumes of data are aggregated into a single, unified repository to optimize a global objective function. This centralization allows for the direct processing of diverse data sources, ensuring the model achieves a comprehensive understanding of the feature space.⁹ In stark contrast, federated learning fundamentally distributes the optimization process, bringing the computational algorithm to the data rather than transporting the data to a centralized compute cluster.⁶

The primary mathematical objective of a federated learning system is to train a global model, defined by a parameter vector $\mathbf{X} \in \mathbb{R}^d$, where d represents the total dimensions of the model. This training occurs across multiple communication rounds utilizing distributed datasets that reside strictly on K different participating clients.¹⁰ The global optimization problem is typically formulated to minimize an aggregated objective function $F(\mathbf{X})$, which represents the weighted sum of the local empirical risks calculated across all participating edge devices.¹⁰

During a standard federated communication round, a central coordinating server orchestrates the learning process by broadcasting the current state of the global model, X_t , to a selected subset of available clients. Each participating client k initializes its local model with X_t and executes a local training routine—most commonly utilizing variations of Stochastic Gradient Descent (SGD)—over its private, localized facial dataset D_k .¹¹ Following the completion of the local training epochs, the client computes a model update, often expressed as a gradient ΔX_k , and transmits solely this differential back to the central server.¹¹ The server then aggregates these inbound updates, frequently utilizing a weighted averaging mechanism such as Federated Averaging (FedAvg), to compute the updated global model X_{t+1} .¹¹ This updated model is subsequently relayed back to the edge devices, initiating the next training epoch.⁶

Experimental findings utilizing established datasets, such as the CelebA dataset, have demonstrated that employing federated learning in both supervised and unsupervised face recognition systems yields dual benefits. Primarily, it safeguards privacy since the original facial data remains confined to the edge devices.⁶ Furthermore, the empirical results confirm that the aggregated global model is capable of achieving performance metrics that are nearly identical to those produced by individual, centrally trained models, validating the efficacy of the federated approach.⁶ However, the assumption that this structural separation of data and computation provides absolute privacy is fundamentally flawed. The transmission of millions of updated parameters inherently encodes the statistical properties and intricate feature representations of the local data, exposing the system to highly sophisticated inference and reconstruction attacks.⁸

The Threat Landscape Part I: Breaches of Confidentiality and Model Inversion

The decentralized and collaborative nature of federated learning drastically expands the potential attack surface. Adversaries—ranging from external hackers intercepting communication channels to malicious insiders, compromised edge devices, or even an "honest-but-curious" central server—can exploit the shared gradients to compromise the confidentiality of the biometric data.

Gradient and Model Inversion Attacks

Model Inversion (MI) attacks represent one of the most severe threats to the privacy guarantees of federated learning. These attacks exploit the fundamental reality that a machine learning model inherently "knows" and retains information about the data upon which it was trained.¹² Gradient inversion attacks are a specialized and highly potent subset of MI specifically tailored for federated environments. In this scenario, an adversary attempts to reconstruct the

original, raw input data (the facial images) by capturing and exploiting the gradient information shared during the model update process.⁸ Because the local gradient $\nabla f(X)$ is computed directly using the client's private facial data, the gradient vector is mathematically tethered to the input space.⁸ By iteratively adjusting a synthetic, randomized "dummy" input and its corresponding dummy label until the gradients produced by this dummy data precisely match the intercepted gradients from the victim client, adversaries can successfully recover the actual training data, leading to catastrophic privacy leakage.⁸

The efficacy and real-world viability of these attacks have been starkly demonstrated through extensive empirical research. In foundational studies (Fredrikson et al., 2015), researchers developed the first model inversion attack algorithm explicitly targeting facial recognition systems. They demonstrated that attackers could produce highly recognizable images of individuals' faces using nothing more than basic API access to the target system and the name of the intended victim.¹³ Subsequent, more advanced iterations of this attack methodology, such as those executed by researchers at Carnegie Mellon University, revealed even more alarming vulnerabilities. Their experiments demonstrated the ability to reconstruct actual human faces from a facial recognition model without ever having access to the original photographs.¹² The attack sequence initiated with garbage input—specifically random pixel noise—which was systematically fed into the facial recognition API. The API returned a confidence score indicating the degree of similarity between the random noise and a specific target individual. By utilizing these confidence scores as an optimization metric, the attackers iteratively refined the pixel noise until a clear, identifiable facial image emerged.¹²

In the context of federated learning, once an adversary successfully infers or obtains a private class embedding, they can easily synthesize high-fidelity face images belonging to specific clients.¹⁴ This is frequently achieved by leveraging advanced generative technologies, such as Generative Adversarial Networks (GANs) or algorithmic techniques like DeepInversion, to optimize random noise into photorealistic representations of the target's face.¹⁴

Federated Unlearning Inversion Attacks (FUIA)

A relatively new dimension of the threat landscape involves the exploitation of "federated unlearning" mechanisms. Federated Unlearning (FU) is a process designed to intentionally eliminate the influence of specific, targeted data from a trained model, often to comply with regulatory mandates like the Right-To-Be-Forgotten.¹⁵ However, the process of unlearning itself introduces unique vulnerabilities. The Federated Unlearning Inversion Attack (FUIA) directly undermines the fundamental goals of FU by exploiting the mathematical vulnerabilities inherent in the unlearning process to reconstruct the very data that the system is attempting to forget.¹⁵ By analyzing the delta between the model states before and after the unlearning procedure, adversaries can isolate the specific features associated with the forgotten data, thereby revealing severe flaws in current privacy protection paradigms.¹⁵ Extensive experiments conducted on multiple benchmark datasets across various FU methodologies have

conclusively demonstrated that FUIA effectively exposes the private information of supposedly deleted records.¹⁵

The Threat Landscape Part II: Membership Inference Attacks

While Model Inversion seeks to reconstruct the actual facial image, Membership Inference Attacks (MIA) pursue a different objective: ascertaining whether a specific individual's facial image was utilized within the training dataset of a given model.¹⁶ This form of attack assesses exactly how much a target machine learning model inadvertently reveals about its underlying training data by determining if specific query instances were "members" of the training set.¹⁶

MIAs exploit the mathematical phenomenon of overfitting—the tendency of deep neural networks to exhibit noticeably higher confidence scores, lower entropy, and lower loss when evaluating data they have previously memorized during training, as opposed to novel, unseen data.¹⁷ State-of-the-art MIAs traditionally rely on training hundreds of independent "shadow models" that are architecturally similar but completely independent of the target model. These shadow models are trained on distinct datasets to mimic the behavior of the target model, ultimately generating a massive labeled dataset consisting of "member" and "non-member" confidence vectors.¹⁶ This synthesized data is then utilized to train a sophisticated binary classifier capable of executing the membership inference attack against the actual target system.¹⁶ However, training hundreds of shadow models introduces immense computational overhead. To circumvent this, recent advancements include the Imitative Membership Inference Attack (IMIA), which employs a novel imitative training technique to strategically construct the attack without the need for exhaustive shadow model training.¹⁶

As defenders attempt to mitigate MIA vulnerabilities, adversaries continuously adapt. For instance, if a defender implements a mitigation mechanism that intentionally increases the computational loss of original training data to obfuscate its membership status, sophisticated attackers deploy the "Inverse Membership Inference Attack".¹⁹ This adaptive adversary learns the specific mechanism of the defense and inversely classifies data exhibiting abnormally high loss as the definite member, effectively bypassing the intended security measure.¹⁹ Conversely, external adversaries utilizing attacks like Ob-MALT may simultaneously classify data with high loss as non-members and low loss as members, requiring defenders to balance multiple, conflicting threat vectors simultaneously.¹⁹

Empirical investigations comparing MIA methodologies across both federated and centralized learning environments reveal that information leakage can be significantly mitigated through a combination of knowledge distillation and conventional regularization techniques.¹⁷ Strategies such as Gaussian dropout, the strategic injection of Gaussian noise, and activity regularization have proven effective at masking the confidence disparities that MIAs rely upon, thereby

reducing the risk of membership disclosure in federated settings.¹⁷

The Threat Landscape Part III: Byzantine Failures and Model Integrity

While Model Inversion and Membership Inference focus entirely on breaching confidentiality, a separate class of attacks targets the fundamental integrity of the facial recognition model itself. These integrity-compromising actions are broadly categorized as Byzantine failures and manifest primarily through data poisoning and model poisoning attacks.⁵

Data Poisoning Attacks

Data poisoning is a sub-branch of poisoning attacks executed at the very inception of the training phase.²¹ In a federated learning architecture, the central server relies entirely on the premise that edge devices are training on valid data. This trust grants malicious end-users or compromised devices full authority over their local training environment. In a data poisoning attack, an adversary intentionally alters the local training dataset. This can be achieved by injecting imperceptible noise into the facial images or, more commonly, by executing a label-flipping attack where the labels of samples are deliberately swapped from one class to another (e.g., labeling images of Person A as Person B).²⁰

Because the central server cannot inspect the localized raw data, it blindly accepts the model updates derived from these corrupted datasets. Consequently, the global model incorporates mislabeled data representations, significantly increasing the probability that the system will incorrectly classify samples during the testing or inference phase.²¹ While data poisoning is detrimental, its impact is often diluted during the global aggregation phase, especially if the malicious client represents only a small fraction of the total participating network.²²

Model Poisoning Attacks

Model poisoning is a significantly more direct, sophisticated, and destructive threat. In this scenario, the adversary bypasses the manipulation of raw data entirely. Instead, the purpose of a model poisoning attack is to arbitrarily and maliciously manipulate the actual model updates (the gradients) before they are transmitted back to the central server.²⁰ By directly corrupting the mathematical parameters, attackers can cause the global model to rapidly deviate from its normal optimization trajectory. This results in either severely degraded overall model performance or, more insidiously, the successful installation of hidden "backdoors" into the final global model.²⁰ A backdoored facial recognition model might function perfectly under normal conditions but will consistently misclassify a specific individual or bypass authentication when presented with a specific, attacker-defined trigger.²⁰

Empirical studies have systematically evaluated the vulnerability of federated learning to local model poisoning. Results spanning multiple real-world datasets indicate that local model

poisoning attacks can inflict devastating damage, resulting in 1.5 to 60 times higher reductions in the overall accuracy of FL models when directly compared to previously discovered data poisoning techniques.²⁴ These findings demonstrate that many existing federated learning algorithms—even those explicitly claimed to be robust against Byzantine failures—are significantly more susceptible to direct model parameter manipulation than previously theorized.²³

Attack Category	Primary Target	Mechanism of Exploitation	Severity & Impact on Biometric Systems
Gradient Inversion	Data Confidentiality	Iterative mathematical matching of dummy data gradients to intercepted client updates.	Critical: Capable of reconstructing raw, identifiable facial images, resulting in total privacy loss.
Membership Inference	Data Confidentiality	Statistical analysis of confidence score and loss disparities between training and unseen data.	High: Violates user anonymity and breaches strict GDPR compliance regarding data presence.
Data Poisoning	Model Integrity	Deliberate mislabelling or the addition of adversarial noise to localized facial datasets prior to training.	Moderate to High: Degrades recognition accuracy for specific demographic classes or individuals.
Model Poisoning	Model Integrity	Direct, arbitrary manipulation of the gradient vector transmitted to the	Critical: Can install hidden backdoors or completely destroy global model

		central server.	convergence and accuracy.
Federated Unlearning Inversion	Data Confidentiality	Exploiting the delta between pre- and post-unlearning model states to recover targeted features.	High: Nullifies the legal and functional benefits of the Right-To-Be-Forgotten.

Cryptographic Defenses and Secure Multiparty Computation

To counteract these highly sophisticated threat vectors, the research community has engineered an extensive array of cryptographic, statistical, and algorithmic defenses designed specifically to fortify the federated facial recognition pipeline without completely sacrificing model utility.

Differential Privacy (DP)

Differential Privacy serves as a foundational statistical defense against both Model Inversion and Membership Inference attacks. The core mechanism of DP involves the intentional injection of calibrated cryptographic noise—typically drawn from Gaussian or Laplacian distributions—into the local gradients before they are transmitted to the central server.³ By perturbing the updates, DP ensures that the final output of the federated learning algorithm is not overly dependent on any single facial image or data point within the dataset. This statistical obfuscation masks the presence or absence of specific individuals, directly neutralizing the confidence disparities exploited by MIAs.³

However, the application of Differential Privacy in the precise domain of facial recognition introduces a highly sensitive privacy-utility tradeoff. Excessive noise injection irrevocably degrades the precision of the facial feature extractors, rendering the model useless for high-security authentication. Conversely, insufficient noise leaves the model mathematically vulnerable to gradient inversion.⁴ Recent empirical assessments demonstrate that defense systems integrating DP alongside adaptive noise injection, hybrid cryptographic encryption, and ensemble-based obfuscation can mitigate up to 50 percent of model inversion attacks relative to baseline models without driving F1 accuracy scores below the critical threshold of 0.85.⁴ Furthermore, transforming these mathematical defenses to align with essential global

data privacy laws, such as GDPR and HIPAA, provides vital guidelines for the ethical and regulation-sensitive deployment of privacy-preserving machine learning models.⁴

Homomorphic Encryption (HE) and Secure Multiparty Computation (SMPC)

While Differential Privacy provides statistical guarantees, Homomorphic Encryption (HE) and Secure Multiparty Computation (SMPC) provide absolute cryptographic guarantees.²⁶ HE is a revolutionary cryptographic technique that allows the central server to mathematically aggregate gradient updates while the data remains entirely encrypted in ciphertext.²⁵ This ensures that the central server—even if compromised or acting in an "honest-but-curious" manner—never observes the raw, unencrypted model updates, thereby eliminating the risk of server-side gradient inversion.²⁵ Advanced implementations in biometric systems have successfully utilized the CKKS algorithm from the SEAL library to encrypt normalized face feature vectors, enabling secure operations without decryption.¹¹ Additionally, to save computation costs associated with the encryption of query face images, novel multi-matching schemes have been proposed to match an encrypted face query directly against clustered faces residing in a secure repository.¹¹

Complementary to HE, Secure Multiparty Computation (SMPC) protocols empower a network of inherently distrustful devices to collaboratively calculate an aggregate value without disclosing their individual private data.²⁸ In a typical SMPC setup for federated learning, a set of devices individually possess private values (their model updates). Through cryptographic collaboration, they compute the required sum of the updates while mathematically ensuring that no single device—nor the central server—discloses or intercepts any information beyond what can be inferred from the final aggregated result.¹¹

Despite their robust security profiles, both HE and SMPC introduce immense computational complexity, massive memory requirements, and significant communication bandwidth overheads. These resource demands are often prohibitive for standard edge-device facial recognition applications, necessitating careful scheduling policies to minimize the number of active users required per round to alleviate throughput restrictions.²⁸

Algorithmic Defenses: Byzantine-Robust Aggregation Strategies

To defend against the insidious threats of data and model poisoning, the standard Federated Averaging (FedAvg) protocol must be replaced with robust, resilient aggregation strategies. The central challenge lies in developing rules that can identify and discard malicious updates without inadvertently discarding benign updates from clients with highly distinct, non-IID data distributions.⁵

The FLRAM Architecture

The Federated Learning Robust Aggregation Method (FLRAM) represents a leading active defense strategy explicitly designed to counter Byzantine poisoning attacks before the framework is compromised.³¹ FLRAM departs from basic defenses that rely on simple gradient median estimation or rudimentary clipping transformations.³¹ Instead, FLRAM is structured around three core operational stages: Anomalous Gradient Detection, Credibility Assessment, and Global Gradient Aggregation.³¹

During the Anomalous Gradient Detection phase, the technique rigorously scrutinizes both the magnitudes and the specific signs of the gradients uploaded by the participating clients, operating on the premise that both data and model poisoning will manifest as detectable geometric deviations.³¹ FLRAM utilizes sophisticated machine learning anomaly detection techniques, specifically integrating Isolation Forests and an improved Density-Based Spatial Clustering of Applications with Noise (DBSCAN) algorithm. The improved DBSCAN is specifically tailored for clustering gradient signs, significantly enhancing the accuracy of abnormal gradient identification.³¹

Following detection, the server constructs a dynamic Credibility Matrix to assess the ongoing trustworthiness of each client. During the Global Gradient Aggregation phase, the central server actively filters out clients identified as abnormal and selectively aggregates model updates only from benign clients demonstrating high credibility scores.³¹ Theoretical proofs confirm the convergence properties of the FLRAM algorithm, assuming a strictly controlled learning rate.³¹ Experimental validations on benchmark datasets demonstrate that FLRAM maintains superior robustness and high model accuracy even when confronting sophisticated attack patterns like Min-Max, Fang, and standard Byzantine attacks, remaining stable as long as the proportion of malicious clients does not exceed 50 percent of the total network participants.³¹

Divide-and-Conquer and cGAN Validation Frameworks

Alternative robust aggregation approaches include the Divide-and-Conquer (DnC) algorithm. Motivated by the extreme vulnerability of standard FL to local model poisoning, DnC actively isolates the principal components of malicious updates.²⁴ Empirical evaluations demonstrate that DnC significantly outperforms all existing Byzantine-robust algorithms in defeating model poisoning, proving to be 2.5 to 12 times more resilient across diverse datasets and model architectures.²⁴

Furthermore, cutting-edge defense frameworks are beginning to leverage advanced generative AI to ensure update integrity. A major limitation of many detection-based methods is their reliance on rigid heuristics or the prerequisite of a trusted, external dataset for validation scoring.⁵ To resolve this, researchers have developed frameworks that deploy a conditional Generative Adversarial Network (cGAN) directly at the central server. The cGAN synthesizes

highly representative data dynamically, which is then utilized to validate the incoming client updates.⁵ This elegant approach eliminates the dangerous reliance on external trusted datasets, adapts fluidly to diverse and evolving attack strategies, and integrates seamlessly into standard federated workflows, successfully distinguishing malicious from benign clients while preserving overall model accuracy.⁵

Overcoming Non-IID Data Heterogeneity in Biometric Distributions

While advanced cryptographic protocols and Byzantine-robust aggregation secure the infrastructure, the most persistent and mathematically complex barrier to achieving high-accuracy federated face recognition is data heterogeneity, strictly categorized as non-independent and identically distributed (non-IID) data.⁹

In idealized machine learning environments, data is assumed to be evenly distributed and standardized. However, in real-world federated deployments, the facial data captured by discrete clients—such as high-end smartphones, low-resolution IoT security cameras, and varying geographic locations—exhibits massive statistical variance.² These variations manifest across multiple axes, including severe demographic skew, wildly divergent lighting intensities, diverse facial expressions, aging effects, occlusions, and varied hardware sensor quality.² Despite the critical significance of non-IID data in federated learning, there remains a notable lack of consensus among researchers regarding its precise classification and quantification, prompting systematic reviews to establish detailed taxonomies for non-IID data, partition protocols, and mathematical metrics to quantify data heterogeneity.⁹

The consequence of this heterogeneity is a phenomenon known as "client drift." Because each client is training locally on a highly skewed, non-representative slice of the overall data distribution, the local optimum determined by the client will diverge significantly from the true global optimum required by the central model.⁹ When these divergent gradients are averaged together, the resulting global model experiences severe degradation in recognition accuracy, excruciatingly slow training times, and prolonged communication rounds to reach convergence.⁹ Comprehensive data partitioning experiments evaluating state-of-the-art FL algorithms consistently confirm that non-IID distributions bring significant challenges, and no single existing algorithm universally outperforms others across all highly heterogeneous cases.³²

Addressing the non-IID challenge in facial recognition requires systemic algorithmic interventions. Current state-of-the-art solutions fall into several primary methodological categories:

1. **Client Drift Mitigation and Aggregation Schemes:** These techniques involve modifying the local objective functions of the clients to constrain their updates, ensuring that the

local model remains tethered to the global trajectory to a certain degree.³³ Advanced aggregation schemes improve the model fusion mechanism at the server to better account for variance, while data sharing techniques introduce limited public datasets or synthesized data to construct a more balanced statistical distribution.³³ The efficacy of these methods in highly sensitive medical imaging—such as brain tumor detection using MRI scans—demonstrates their potential. In studies simulating real-world non-IID conditions across multiple clients, models like VGG19 utilizing advanced algorithms such as FedProx and Scaffold achieved high test accuracies of 98.24% and 98.45% respectively, outperforming standard FedAvg and showcasing FL's capability to handle distributed, heterogeneous data without violating privacy.³⁴

2. **Self-Supervised and Transfer Learning Frameworks:** Handling the difficulty of collecting and labeling diverse samples for countless real-world scenarios is profoundly difficult. One solution involves transfer learning: first learning highly generic models centrally, and then transferring them to specific, constrained application scenarios.² Alternatively, frameworks like FedAffect leverage self-supervised learning to combat small sample sizes and unlabeled private data.² FedAffect isolates and updates a network of feature extractors on unlabeled private facial data to learn highly diverse, robust facial representations, which intrinsically helps mitigate the impact of heterogeneous data distributions across the network.²
3. **Advanced Sensing and Geometric Data:** Hardware solutions are also being explored. The integration of 3D sensor technology provides crucial geometric facial information that remains highly invariant to unfavorable acquisition conditions, such as severe lighting or pose changes, significantly improving recognition accuracy over traditional 2D techniques.² However, the current scarcity of comprehensive 3D face recognition databases hinders rapid deep learning development in this specific area.²

State-of-the-Art Facial Recognition Frameworks (2024-2026)

The period spanning 2024 to 2026 has witnessed the introduction of specialized federated learning frameworks explicitly engineered to resolve the unique paradoxes of facial recognition: maximizing feature distinctiveness between millions of identities while maintaining strict, decentralized privacy.⁷

FedGC: Gradient-Corrected Joint Averaging

Traditional, centralized facial recognition systems rely heavily on deep neural networks utilizing a shared, fully connected classification layer—often employing specific loss functions like ArcFace or CosFace.¹⁴ This shared layer contains the precise class centers, or "embeddings," for every known identity in the network. In a federated setting, sharing these specific class embeddings among clients fundamentally violates the core premise of privacy, as these high-fidelity embeddings can be easily extracted and reverse-engineered via Generative

Adversarial Networks (GANs) or techniques like DeepInversion to synthesize photorealistic images of the clients.¹⁴ Conversely, completely withholding these embeddings prevents the network from properly calculating standard softmax loss, leading to network collapse or devastatingly poor model performance, particularly when applied to deep network structures.¹⁴

To resolve this complex mathematical paradox, researchers introduced Federated Learning with Gradient Correction (FedGC).² Developed by Niu and Deng, FedGC is a privacy-preserving co-learning framework that introduces a novel training strategy tailored specifically for privately dispersed, non-IID face data.² The defining innovation of FedGC is that it guarantees each client securely holds and maintains its own private, fully available class embeddings locally.²

To ensure that the network updates follow the optimal optimization trajectory without sharing the embeddings, FedGC addresses the local optimization problem by correcting the gradients entirely from the perspective of backward propagation.² FedGC integrates a specialized softmax-based regularizer designed to correct the gradients of the class embeddings.² By precisely calculating and injecting a unique "cross-client gradient term," FedGC combines local optimization for softmax regularization with robust inter-client optimization, ensuring that the entire network updates in the standard softmax direction.² Theoretical mathematical proofs confirm that the FedGC gradient injection constitutes a valid loss function structurally identical to standard softmax, seamlessly solving the privacy-preserving limitations inherent in traditional FL methods.⁷ Extensive empirical experiments demonstrate that FedGC achieves performance virtually identical to conventional centralized methods utilizing full training datasets, while entirely eliminating the risk of embedding-based privacy leakage.⁷

VOIDFace: Patch-Based Security and User Sovereignty

Emerging as a paradigm-shifting architecture in 2025, the VOIDFace system, developed by researchers at the VIS Team, integrates multi-network processing with unprecedented user sovereignty mechanisms.¹ VOIDFace addresses the dual, intertwined threats of sophisticated Model Inversion attacks and the systemic issue of unauthorized training data replication through a novel "patch-based" distributed training mechanism utilizing cryptographic secret sharing.³¹

Unlike traditional systems that process entire facial inputs, VOIDFace fragments the biometric data. During the initial training data preparation phase, an individual facial patch (denoted as P_1) is mathematically encrypted into two distinct random grids utilizing a random number generator function.³⁹ One of these random grids operates as the private share (PS_1) for the patch, while the other functions as the critical authentication share (AS).³⁹ When a training request is initiated, a trusted third party must first verify user authorization by checking the

authentication shares. Only upon successful validation are the training request and authentication shares dispatched to N_p independent, strictly non-communicating selected training workstations.³⁹ Each workstation collects the private share from its relevant counterpart and subsequently reconstructs the mathematical representations of the patches necessary for training.³⁹ Because the training occurs utilizing secret shares distributed across independent workstations rather than whole, unencrypted facial images, Model Inversion attacks attempting to reconstruct the source face become mathematically intractable, ensuring high-level security.³⁹

Furthermore, VOIDFace pioneers the highly sought-after technical implementation of the Right-To-Be-Forgotten (RTBF) property within biometric AI architectures.³⁹ Because the individual's data is fundamentally structured through cryptographic shares tied directly to an authentication protocol, a user can instantly and unilaterally revoke access. The moment the user exercises their RTBF, the authentication share is nullified. Consequently, the distributed workstations can no longer authorize or reconstruct the mathematical representations, organically and instantly excising the user's data from all future training iterations without requiring a catastrophic retrain of the entire global model.³⁹ Experimental evaluations conducted on extensive databases, such as the VGGFace2 dataset, demonstrate that VOIDFace provides verified RTBF capabilities, improved data control, and formidable security while maintaining highly competitive facial recognition performance and operational efficiency.³⁹

Specialized Contextual Learning Frameworks

Beyond the generalized recognition capabilities of FedGC and VOIDFace, the academic ecosystem encompasses highly specialized frameworks designed specifically to tackle extreme data sparsity and localized contextual challenges:

- **FedFace:** Engineered specifically for extreme collaborative learning environments where an individual mobile device might possess only a single recognized image of the user, presenting a severe few-shot learning challenge.² FedFace utilizes an extended regularizer to mathematically ensure that class embeddings are properly separated within the highly constrained feature space. Furthermore, it initializes these embeddings utilizing a novel average feature initialization scheme, successfully overcoming the devastating lack of local data variety.²
- **PrivacyFace:** This framework tackles the non-IID data divergence problem by utilizing a Differential Private Local Clustering (DPLC) mechanism.² Instead of broadcasting highly sensitive raw gradients, PrivacyFace purifies the class information locally. It then broadcasts this differentially private class information to the global server, facilitating precise client alignment through a specialized "consensus-aware loss" mechanism, bridging the gap between highly heterogeneous clients without compromising privacy.²
- **Privacy Projector for AIoT:** Designed for industrial Artificial Intelligence of Things (AIoT)

deployments, this framework splits the underlying models into a public module (where only necessary gradients are shared with the server) and a localized private module. The private module acts as a "projector" to evaluate feature quality locally on the edge device, drastically expediting training convergence times in resource-constrained environments.²

Framework Architecture	Primary Algorithmic Innovation	Key Biometric Challenge Addressed	Core Operational Mechanism
FedGC (Gradient Correction)	Cross-client gradient injection	Resolving the lack of a shared classification layer	Employs a softmax-based regularizer to correct backpropagation trajectories securely.
VOIDFace	Right-To-Be-Forgotten (RTBF) capability	Combating Model Inversion and data replication	Utilizes patch-based cryptographic secret sharing across non-communicating networks.
FedFace	Extended regularizers & average feature initialization	Managing extreme few-shot learning scenarios	Optimizes separation of embeddings when clients possess highly limited local imagery.
PrivacyFace	Consensus-aware loss function	Bridging non-IID data divergence across clients	Broadcasts purified data via Differential Private Local Clustering (DPLC).
FedAffect	Self-supervised facial representation	Addressing vast facial expression	Isolates and updates feature extractors continuously on

		variability	unlabeled local data streams.
--	--	-------------	-------------------------------

Hardware Constraints, Resource Exhaustion, and Edge Device Limitations

The theoretical elegance of federated learning is frequently compromised by the harsh physical realities of deployment. FL mandates that a multitude of distributed devices frequently exchange massive deep neural network (DNN) model parameters over varied networks.⁴³ For state-of-the-art facial recognition models—which frequently contain millions to billions of parameters—this introduces a critical and often paralyzing communication overhead.⁴³ In real-world practical deployments, federated systems consist of highly heterogeneous devices ranging from resource-rich flagship smartphones to latency-sensitive edge devices and severely battery-constrained IoT sensors.⁴⁴ These edge devices inherently lack the computational power, massive memory, and thermal dissipation capabilities found in centralized High-Performance Computing (HPC) cloud clusters.⁴⁵

The Evolution of Hardware Accelerators: GPUs, FPGAs, and ASICs

The execution of highly complex, iterative deep learning training epochs locally on edge nodes necessitates specialized hardware accelerators. Relying on traditional, lightweight CPUs is entirely impracticable due to low throughput, while standard high-performance GPUs are excessively power-hungry, bulky, and thermally restricted for compact edge deployments.⁴⁶ Consequently, the industry is increasingly investigating and adopting Field-Programmable Gate Arrays (FPGAs) and Application-Specific Integrated Circuits (ASICs) for edge-based deep learning.⁴⁶

FPGAs offer a highly attractive proposition due to their exceptional power efficiency. Because FPGAs allow the physical hardware logic to be fundamentally reconfigured to optimize specific machine learning algorithms, they can be drastically more efficient; research initiatives like Microsoft's Catapult Project estimate that FPGAs are up to 10 times more power-efficient than traditional GPUs.⁴⁶ Notable implementations, such as Microsoft's Project Brainwave, utilize soft DNN engines rendered with Intel's Stratix FPGAs to deliver real-time, low-latency artificial intelligence.⁴⁶ However, FPGAs present significant drawbacks for dynamic federated learning environments. They suffer from lower clock frequencies (typically capped around 300 MHz, which is four to five times slower than typical ASICs) and require highly complex programmable logic, making the deployment of continuously updating, massive FL models challenging.⁴⁶

Conversely, ASICs deliver vastly superior deep learning efficiency and programmability. While

they possess lower reconfigurability than FPGAs, ASICs introduce complex logic architectures specifically tailored for deep learning operations, resulting in much higher overall efficiency and operating frequencies.⁴⁶ Crucially, because federated learning demands substantial local storage for the continuous processing of local facial datasets, the reduced network overhead and highly optimized off-chip memory access characteristics of ASICs increasingly render them the superior architectural choice for enabling sophisticated FL implementation requirements in edge devices.⁴⁶

Communication Efficiency: Model Compression, Quantization, and Knowledge Distillation

To mitigate the crippling transmission bottlenecks and reduce the massive convergence times associated with transmitting full-precision model updates, researchers have pioneered highly advanced model compression and quantization techniques.

Quantization structurally reduces the numerical precision of the weights and gradients being transmitted—for example, shifting from massive 32-bit floating-point parameters to highly compact 8-bit or even ternary representations—drastically shrinking the payload volume and subsequent transmission delay.⁴³ Product Quantization (PQ), a compression technique originally developed for nearest-neighbor search, has been successfully reformulated and applied to accurately represent and compress model updates specifically within the massive linear layers of deep neural networks.⁴⁹

A highly notable and sophisticated advancement in this domain is the FedDQ framework. FedDQ implements an innovative "descending-trend" quantization scheme.¹⁰ Theoretical mathematical convergence proofs for non-convex scenarios establish that by allowing each participating client to individually determine its optimal quantization bit-depth based on the optimization trajectory—specifically by setting the learning stepsize as $\eta = 1/L\sqrt{K\tau}$ to satisfy the first-order stationary condition—FedDQ drastically reduces the overall communication volume compared to traditional ascending schemes.¹⁰ Paradoxically, this optimized compression scheme allows the model to achieve faster global convergence by consuming significantly fewer communication rounds.¹⁰

Other leading approaches leverage advanced Knowledge Distillation at the edge, exemplified by the FedDT framework. Under this paradigm, a massive, highly accurate "teacher" model trains a heavily compressed "student" model locally on the client device. Only the highly quantized parameters of the student model are transmitted to the server. By capitalizing on the mathematical properties of knowledge distillation, the student model successfully conceals the underlying heterogeneity of the local network, serving to mitigate both device-level hardware limitations and the statistical heterogeneity of the non-IID data.⁴⁷ Comprehensive evaluations utilizing benchmark datasets like MNIST and CIFAR-10 confirm that FedDT attains up to 7.85% higher model accuracy while systematically reducing communication overhead by an

astounding average of 78% compared to baseline methods, providing a robust, lightweight solution for bandwidth-constrained FL systems.⁴⁷

The FTTE System Architecture for Extreme Scalability

Synthesizing these myriad efficiencies, the FTTE framework represents a breakthrough in practical, scalable, and resource-efficient FL systems designed explicitly for real-world federated networks dominated by resource-constrained edge devices.⁵⁰ Unlike traditional systems that struggle beyond a few dozen participants, FTTE is engineered to scale up to 500 simultaneous clients while remaining highly robust under conditions heavily populated by "stragglers"—clients hampered by severe computational delays and massive non-IID data distribution skews.⁵⁰

Empirical benchmarks demonstrate the staggering superiority of the FTTE architecture. When evaluated on complex image datasets such as CIFAR-10, Oxford-IIIT Pet, and Flowers-102, FTTE achieved significantly higher target accuracies (e.g., 71.9% on CIFAR-10) compared to legacy models.⁵⁰ More importantly, FTTE achieves this accuracy while converging on average 81 percent faster, consuming 80 percent less on-device memory footprint, and requiring a 69 percent reduction in overall communication payload size when directly compared to traditional synchronous Federated Averaging (FedAvg/SyncFL).⁵⁰ These metrics establish FTTE as a premier, highly practical solution for massive-scale real-world FL deployments on highly heterogeneous, battery-constrained edge networks.⁵¹

Optimization Strategy	Core Mechanism	Primary Impact Matrix	Target Limitation Addressed
FedDQ (Quantization)	Descending-trend bit reduction with dynamic stepsize (η).	Reduces communication volume while accelerating convergence rounds.	Massive payload bandwidth exhaustion.
FedDT (Distillation)	Local teacher models train ternary quantized student models for upload.	Decreases overhead by 78% and improves accuracy by masking heterogeneity.	Device constraints and non-IID data skew.

Product Quantization (PQ)	Applies nearest-neighbor compression to linear layer updates.	Highly efficient mathematical compression of dense parameter matrices.	Inefficient parameter transmission.
FTTE Framework	Optimized resource-efficient pipeline scaling to 500 clients.	Cuts memory usage by 80%, payload by 69%, and speeds convergence by 81%.	On-device RAM, battery drain, and straggler delays.
ASIC Acceleration	Application-specific off-chip memory access and complex logic.	Vastly outperforms FPGAs in clock frequency and DL efficiency for local training.	Edge processor thermal and computational limits.

Future Horizons, Generative AI, and Multimodal Integration

As the domain of federated facial recognition rapidly approaches widespread commercial and industrial maturity, several distinct, highly interdisciplinary research trajectories are emerging that will undoubtedly define the next generation of decentralized biometric architectures.

A primary focus for the immediate future—heavily emphasized in preeminent academic forums such as the Computer Vision and Pattern Recognition Conference (CVPR) through events like the 2026 FedVision workshop and multiple synthetic data symposiums (SynData4CV, NeXD)—is the integration of advanced Generative Artificial Intelligence into the federated pipeline.⁵²

Researchers are actively exploring how generative image modeling can serve as a secure, privacy-compliant data source for training computer vision models. By utilizing generative AI to locally synthesize vast quantities of highly diverse, statistically representative—yet entirely artificial—facial variations at the edge, frameworks can pre-train local models to overcome data sparsity without ever collecting actual human images.⁵² Furthermore, the application of generative models, such as conditional GANs, directly on the central server to synthesize representative data for validating client updates eliminates the necessity of maintaining trusted, external validation datasets, significantly enhancing the autonomy and security of the

aggregation process.⁵

Secondly, the integration of multi-modal biometrics is recognized as a critical requirement to offset the inherent limitations and spoofing vulnerabilities of pure optical facial recognition. Modern federated approaches are beginning to fuse traditional visual facial embeddings with concurrent, continuous physiological signals. A prominent example involves combining localized facial feature extraction with optical Photoplethysmography (PPG) measurements and electrocardiogram (ECG) traces seamlessly acquired via the rapidly expanding Internet of Healthcare Things (IoHT) ecosystem.³ This sophisticated multimodal fusion not only dramatically increases overall authentication accuracy but intrinsically integrates continuous "liveness" detection, effectively neutralizing physical presentation attacks—such as holding a high-resolution photograph or screen to a camera—without requiring any additional raw data transmission to the central server.³

Furthermore, the development of highly adaptive, continuous learning algorithms is paramount.²⁶ Human facial data is not static; it is highly subject to gradual "concept drift" as individuals age, alter their appearance, or change their environments. Future federated learning algorithms must possess the sophisticated capability to dynamically react to these temporal variations in data distribution and shifts in edge device activity.²⁷ This adaptability will preserve the relevance and high accuracy of the localized biometric models over extended timeframes, actively preventing the catastrophic forgetting of previously learned identity embeddings while seamlessly incorporating new data.²⁷

Finally, the academic community faces a pressing imperative to formalize and standardize the taxonomy and quantification of non-IID data. Currently, the lack of universally accepted mathematical metrics to precisely quantify data heterogeneity across clients severely stymies comparative analysis between competing FL frameworks.⁹ Establishing standardized partition protocols and rigorous non-IID metrics will enable significantly more precise engineering of client drift mitigation strategies.⁹ Concurrently, as global artificial intelligence legislation tightens, the seamless, verifiable translation of deeply technical mathematical defenses—such as dynamic noise injection, adaptive differential privacy, and homomorphic encryption—into auditable compliance mechanisms for GDPR, HIPAA, and emerging AI acts will be absolutely vital for the legal, ethical deployment of these systems in enterprise, government, and healthcare environments.⁴

Conclusion

The complex intersection of deep learning facial recognition and federated learning architectures represents one of the most critical, mathematically intricate, and rapidly evolving subfields within contemporary artificial intelligence research. By fundamentally decoupling the necessity of massive data centralization from the sophisticated process of algorithmic optimization, federated learning offers a structurally superior, highly scalable paradigm for the

rigorous protection of sensitive, immutable biometric identifiers.

However, as this exhaustive analysis demonstrates, the transition to a decentralized architecture is not a panacea; it is not immune to highly sophisticated exploitation. Inherent vulnerabilities, such as gradient-based Model Inversion, advanced Imitative Membership Inference, and devastating Byzantine model poisoning attacks, demand the rigorous, continuous application of advanced cryptographic protocols, anomaly detection mechanisms, and highly robust aggregation methodologies like FLRAM, DnC, and cGAN-based validation. Simultaneously, the persistent mathematical challenges introduced by severe non-IID data distributions across heterogeneous client networks require profound structural algorithmic innovations, beautifully exemplified by the FedGC framework's cross-client gradient injection and the integration of highly localized self-supervised learning techniques.

Moreover, the physical limitations of edge computing hardware heavily dictate the ultimate feasibility of real-world deployment. The necessary transition toward highly efficient ASIC hardware accelerators, coupled with aggressive, mathematically proven model quantization schemes like FedDQ and highly optimized network architectures like FTTE, is absolutely critical to overcoming the immense communication and latency bottlenecks inherent in training massive biometric models across distributed global networks. Furthermore, the concept of absolute user sovereignty over biometric data is rapidly shifting from a theoretical legal ideal to an executable technical reality, as demonstrated by the VOIDFace framework's pioneering integration of patch-based cryptographic secret sharing and its functional implementation of the Right-To-Be-Forgotten.

Ultimately, realizing the full protective and operational potential of privacy-focused facial recognition necessitates a highly holistic, deeply interdisciplinary approach. It requires the continuous, seamless synthesis of impenetrable cryptographic security, mathematically rigorous optimization algorithms capable of surviving extreme data heterogeneity, and highly efficient hardware-software co-design. As these disparate disciplines converge, federated learning will solidify its position not merely as a privacy-preserving alternative to centralization, but as the foundational, secure architecture for the next era of ubiquitous, global biometric authentication.

Works cited

1. Federated Learning for Secure and Privacy-Preserving Facial ..., accessed March 20, 2026, <https://visteam.isr.uc.pt/publications/federated-learning-for-secure-and-privacy-preserving-facial-recognition-advances-challenges-and-research-directions/>
2. Federated Learning-Based Face Recognition: Methods, Challenges and Future Prospects - SciTePress, accessed March 20, 2026, <https://www.scitepress.org/Papers/2024/129594/129594.pdf>
3. Authors - The Alan Turing Institute, accessed March 20, 2026, https://www.turing.ac.uk/sites/default/files/2025-01/wp_2_examination_of_privacy

- [_enhancing_technologies_in_biometric_matching_bryan_-_final_0.pdf](#)
4. A Comparative Analysis of Defense Mechanisms Against Model Inversion Attacks on Tabular Data - MDPI, accessed March 20, 2026, <https://www.mdpi.com/2624-800X/5/4/80>
 5. Byzantine-Robust Federated Learning Using Generative Adversarial Networks - arXiv, accessed March 20, 2026, <https://arxiv.org/pdf/2503.20884>
 6. Federated Learning Method for Preserving Privacy in Face ... - arXiv, accessed March 20, 2026, <https://arxiv.org/abs/2403.05344>
 7. The Advancements and Future Prospects of Federated Learning-Based Methods for Biometrics - SciTePress, accessed March 20, 2026, <https://www.scitepress.org/Papers/2024/135253/135253.pdf>
 8. Model Inversion Attacks: A Survey of Approaches and Countermeasures, accessed March 20, 2026, <https://arxiv.org/html/2411.10023v2>
 9. Non-IID data in Federated Learning: A Systematic Review with Taxonomy, Metrics, Methods, Frameworks and Future Directions - arXiv, accessed March 20, 2026, <https://arxiv.org/html/2411.12377v1>
 10. FedDQ: Communication-Efficient Federated Learning with Descending Quantization - arXiv.org, accessed March 20, 2026, <https://arxiv.org/pdf/2110.02291>
 11. Federated Learning Method for Preserving Privacy in Face Recognition System - arXiv, accessed March 20, 2026, <https://arxiv.org/html/2403.05344v1>
 12. Model Inversion Attacks: When AI Reveal Their Secrets - Astra Security, accessed March 20, 2026, <https://www.getastra.com/blog/ai-security/model-inversion-attacks/>
 13. Model Inversion Attacks: Risks and Defenses Explained - SentinelOne, accessed March 20, 2026, <https://www.sentinelone.com/cybersecurity-101/cybersecurity/model-inversion-attacks/>
 14. Federated Learning | PDF | Applied Mathematics - Scribd, accessed March 20, 2026, <https://www.scribd.com/document/738791751/Federated-Learning>
 15. Model Inversion Attack Against Federated Unlearning - ResearchGate, accessed March 20, 2026, [https://www.researchgate.net/publication/400968677_Model_Inversion_Attack_a gainst Federated Unlearning](https://www.researchgate.net/publication/400968677_Model_Inversion_Attack_against_Federated_Unlearning)
 16. Imitative Membership Inference Attack - arXiv.org, accessed March 20, 2026, <https://arxiv.org/html/2509.06796v1>
 17. Comparative Analysis of Membership Inference Attacks in Federated and Centralized Learning - MDPI, accessed March 20, 2026, <https://www.mdpi.com/2078-2489/14/11/620>
 18. Subject Membership Inference Attacks in Federated Learning - Oracle APEX, accessed March 20, 2026, https://apexapps.oracle.com/pls/apex/f?p=LABS:0:103944954854778:APPLICATION_PROCESS=GETDOC_INLINE:::DOC_ID:2920
 19. Fortifying Federated Learning against Membership Inference Attacks via Client-level Input Perturbation - Yinzhi Cao, accessed March 20, 2026,

<https://yinzhicao.org/CIP/DSN-23-CIP.pdf>

20. Threats, attacks and defenses to federated learning: issues, taxonomy and perspectives, accessed March 20, 2026, <https://d-nb.info/1255995327/34>
21. Reducing Defense Vulnerabilities in Federated Learning: A Neuron-Centric Approach, accessed March 20, 2026, <https://www.mdpi.com/2076-3417/15/11/6007>
22. Byzantine-Robust Federated Learning: An Overview With Focus on Developing Sybil-based Attacks to Backdoor Augmented Secure Aggregation Protocols - arXiv, accessed March 20, 2026, <https://arxiv.org/html/2410.22680v1>
23. Local Model Poisoning Attacks to Byzantine-Robust Federated Learning - USENIX, accessed March 20, 2026, <https://www.usenix.org/system/files/sec20-fang.pdf>
24. Manipulating the Byzantine: Optimizing Model Poisoning Attacks and Defenses for Federated Learning - Network and Distributed System Security (NDSS) Symposium, accessed March 20, 2026, https://www.ndss-symposium.org/wp-content/uploads/ndss2021_6C-3_24498_paper.pdf
25. The Global Data Residency Crisis: How Enterprises Can Navigate Geolocation, Storage, and Privacy Compliance Without Sacrificing Performance - Security Boulevard, accessed March 20, 2026, <https://securityboulevard.com/2025/12/the-global-data-residency-crisis-how-enterprises-can-navigate-geolocation-storage-and-privacy-compliance-without-sacrificing-performance/>
26. (PDF) Enhancing healthcare data privacy and interoperability with federated learning, accessed March 20, 2026, https://www.researchgate.net/publication/391567100_Enhancing_healthcare_data_privacy_and_interoperability_with_federated_learning
27. Enhancing healthcare data privacy and interoperability with federated learning - PeerJ, accessed March 20, 2026, <https://peerj.com/articles/cs-2870/>
28. UNSUPERVISED FEDERATED LEARNING FOR FACE RECOGNITION IN DECENTRALIZED ENVIRONMENTS - OpenReview, accessed March 20, 2026, <https://openreview.net/pdf?id=XH3Oihhtvf>
29. Maintaining Privacy in Face Recognition Using Federated Learning Method Woubie, Abraham - HELDA - Helsinki.fi, accessed March 20, 2026, <https://helda.helsinki.fi/bitstreams/25715438-54e0-4e64-a4a3-d63dfe835f83/download>
30. D3.1 — Interim report on Scalable Secure Multiparty Computation, Federated Learning and Unbiased AI techniques and tools, accessed March 20, 2026, https://secured-project.eu/wp-content/uploads/2025/04/D3.1_Interim-report-on-Scalable-Secure-Multiparty-Computation-Federated-Learning-and-Unbiased-AI-techniques-and-tools.pdf
31. FLRAM: Robust Aggregation Technique for Defense against ... - MDPI, accessed March 20, 2026, <https://www.mdpi.com/2079-9292/12/21/4463>
32. Federated Learning on Non-IID Data Silos: An Experimental Study - Flower Baselines 1.27.0, accessed March 20, 2026, https://flower.ai/docs/baselines/niid_bench.html

33. No Fear of Heterogeneity: Classifier Calibration for Federated Learning with Non-IID Data - NIPS, accessed March 20, 2026, <https://papers.neurips.cc/paper/2021/file/2f2b265625d76a6704b08093c652fd79-Paper.pdf>
34. Federated Learning Framework for Brain Tumor Detection Using MRI Images in Non-IID Data Distributions - PMC, accessed March 20, 2026, <https://pmc.ncbi.nlm.nih.gov/articles/PMC12701187/>
35. Towards Federated Learning Driving Technology for Privacy- Preserving Micro-Expression Recognition - IEEE Xplore, accessed March 20, 2026, <https://ieeexplore.ieee.org/iel8/5971803/10979778/10979783.pdf>
36. PFedFace: Personalized Federated Learning for Face Recognition - R Discovery, accessed March 20, 2026, <https://discovery.researcher.life/article/pfedface-personalized-federated-learning-for-face-recognition/c8e4c4d692543e5f9c5d4ffe62151e27>
37. Federated Learning in Computer Vision - Università degli studi di Padova, accessed March 20, 2026, <https://www.research.unipd.it/retrieve/d7f14eaa-02fa-49ca-9de4-9f66498682fc/Federated.pdf>
38. [Quick Review] Federated Learning for Face Recognition with Gradient Correction - Liner, accessed March 20, 2026, <https://liner.com/review/federated-learning-for-face-recognition-with-gradient-correction>
39. VOIDFace: A Privacy-Preserving Multi-Network Face Recognition With Enhanced Security, accessed March 20, 2026, <https://arxiv.org/html/2508.07960v1>
40. VOIDFace: A Privacy-Preserving Multi-Network Face Recognition With Enhanced Security, accessed March 20, 2026, <https://visteam.isr.uc.pt/publications/voidface-towards-an-effective-face-training-data-storage-and-protection-with-right-to-be-forgotten-property/>
41. VOIDFace: Towards an effective face training data storage and protection with Right-To-Be-Forgotten property - ResearchGate, accessed March 20, 2026, https://www.researchgate.net/publication/401530952_VOIDFace_Towards_an_effective_face_training_data_storage_and_protection_with_Right-To-Be-Forgotten_property
42. VOIDFace: A Privacy-Preserving Multi-Network Face Recognition With Enhanced Security, accessed March 20, 2026, https://www.researchgate.net/publication/394439697_VOIDFace_A_Privacy-Preserving_Multi-Network_Face_Recognition_With_Enhanced_Security
43. Communication-efficient federated learning | PNAS, accessed March 20, 2026, <https://www.pnas.org/doi/10.1073/pnas.2024789118>
44. Client Selection in Federated Learning on Resource-Constrained Devices: A Game Theory Approach - MDPI, accessed March 20, 2026, <https://www.mdpi.com/2076-3417/15/13/7556>
45. Efficient Acceleration of Deep Learning Inference on Resource-Constrained Edge Devices: A Review - IEEE Xplore, accessed March 20, 2026, <https://ieeexplore.ieee.org/iel7/5/10015202/09985008.pdf>

46. Federated Learning in Edge Computing: A Systematic Survey - PMC, accessed March 20, 2026, <https://pmc.ncbi.nlm.nih.gov/articles/PMC8780479/>
47. FedDT: A Communication-Efficient Federated Learning via Knowledge Distillation and Ternary Compression - MDPI, accessed March 20, 2026, <https://www.mdpi.com/2079-9292/14/11/2183>
48. Communication-Efficient Federated Learning for Multi-Institutional Medical Image Classification - PMC, accessed March 20, 2026, <https://pmc.ncbi.nlm.nih.gov/articles/PMC9605156/>
49. Reconciling Security and Communication Efficiency in Federated Learning - DIMACS, accessed March 20, 2026, https://www.dimacs.rutgers.edu/~graham/pubs/papers/efficient_fl.pdf
50. FTTE: Federated Learning on Resource-Constrained Edge Devices - arXiv, accessed March 20, 2026, <https://arxiv.org/html/2510.03165v1>
51. (PDF) FTTE: Federated Learning on Resource-Constrained Devices - ResearchGate, accessed March 20, 2026, https://www.researchgate.net/publication/396223599_FTTE_Federated_Learning_on_Resource-Constrained_Devices
52. CVPR 2025 Workshops, accessed March 20, 2026, <https://cvpr.thecvf.com/virtual/2025/events/workshop>
53. CVPR 2025 Unveils the Latest Trends in Artificial Intelligence (AI,) Image and Video Synthesis, Neural Networks, and 3D Computer Vision, accessed March 20, 2026, <https://cvpr.thecvf.com/Conferences/2025/News/Wrap>
54. CVPR 2026 Workshops - The Computer Vision Foundation, accessed March 20, 2026, <https://cvpr.thecvf.com/Conferences/2026/Workshops>